

Project-IV - Parametric/Nonparametric Nonlinear Regression

Md Al Masum Bhuiyan

November 05, 2018

Contents

1	Problem 1: Data Import and Analysis	2
2	Problem 2: Data Partition	3
3	Problem-3: Parametric Nonlinear Models	4
3.1	Problem-3(a) - Model Fitting	4
3.2	Problem-3(b): Hypothesis Testing	5
3.3	Problem-3(c) Fitted curve	6
3.4	Problem-3(d) MSE	6
4	Problem-4: Local Regression Methods	7
4.1	Problem 4(a): KNN regression	7
4.2	Problem 4(b): Kernel Regression	8
4.3	Problem 4(c): Local (Cubic) Polynomial Regression	9
5	Problem 5: Regression/ Smoothing Splines	10
5.1	Problem 5(a): Regression Splines (Natural Cubic)	10
5.2	Problem 5(b): Smoothing Splines	13
6	Problem-6	14

This project is based on a data set jaws.txt, which is concerned about the association between jaw bone length ($y = \text{bone}$) and age in deer ($x = \text{age}$). I am going try out several parametric/nonparametric nonlinear regression models in this low-dimensional ($p = 1$) setting.

1 Problem 1: Data Import and Analysis


```
data <- read.table(file="/Users/masum/Desktop/Fall_2018/Data_Mining/Project/Project_04/jaws.txt")
dat <- data
dim(dat)
```

```
## [1] 54 2
```

```
head(dat)
```

```
##      age      bone
## 1  0.000000  0.00000
## 2  5.112000 20.22000
## 3  1.320000 11.11130
## 4 35.240000 140.65000
## 5  1.632931 26.15218
## 6  2.297635 10.00100
```

```
plot(dat$age, dat$bone, main="Bone vs. Age", xlab="Ages ", ylab="Bones", pch=1, col="orange")
```



Remarks: The association between bones vs. ages (deer jaws) does not look linear.

2 Problem 2: Data Partition

I randomly partition the data into training and test with a ratio 2:1.

```
set.seed(123)
n <- nrow(dat)
split_data <- sample(x=1:2, size = n, replace =TRUE, prob=c(0.67, 0.33))
train <- dat[split_data ==1, ]
test <- dat[split_data ==2, ]
```

Now I indicate the observations that are from train or test obs.

```
dat$partition <- sample(c("train","test"), size=nrow(dat), replace=TRUE, prob=c(.67, .33))
dat
```

##		age	bone	partition
## 1	0.000000	0.00000	train	
## 2	5.112000	20.22000	train	
## 3	1.320000	11.11130	train	
## 4	35.240000	140.65000	test	
## 5	1.632931	26.15218	test	
## 6	2.297635	10.00100	train	
## 7	3.322125	55.85634	train	
## 8	4.043794	46.06754	train	
## 9	5.266018	67.24174	train	
## 10	6.075060	59.63458	train	
## 11	7.236330	60.10914	test	
## 12	8.291007	64.31393	train	
## 13	9.040889	84.92216	test	
## 14	10.881032	92.78716	test	
## 15	11.098745	98.35159	test	
## 16	12.093001	62.22000	train	
## 17	13.998177	85.77307	test	
## 18	14.310943	106.66963	train	
## 19	15.954420	97.90825	test	
## 20	16.076023	105.96718	train	
## 21	17.062700	112.25996	train	
## 22	18.629606	102.48663	train	
## 23	19.330897	114.78824	train	
## 24	20.546239	112.55459	train	
## 25	21.329421	81.11000	train	
## 26	22.264398	123.40000	train	
## 27	23.028779	116.07627	train	
## 28	24.118884	114.72701	train	
## 29	25.314858	105.51583	train	
## 30	26.612978	101.38623	test	
## 31	27.574067	97.77144	train	
## 32	28.130495	118.38879	train	
## 33	29.751232	96.94697	test	

```
## 34 30.123247 113.91214      test
## 35 31.986346  99.48091      test
## 36 32.707306 116.23068    train
## 37 33.726616 106.01806    train
## 38 34.993040 122.96777    train
## 39 35.882134 116.86115    train
## 40 36.829509 122.32905    train
## 41 37.117894  99.07155    train
## 42 38.153487 116.16566    train
## 43 39.439631 123.10436      test
## 44 40.251855 118.56645    train
## 45 41.693744 112.55648    train
## 46 42.649101  99.50362    train
## 47 43.764632 118.69918    train
## 48 44.615721 105.21623    train
## 49 45.855321 108.51228    train
## 50 46.015332 135.00000      test
## 51 47.700156 112.42955    train
## 52 48.964349 101.67611      test
## 53 49.329557 142.00000      test
## 54 50.604097  91.20000    train
```

3 Problem-3: Parametric Nonlinear Models

3.1 Problem-3(a) - Model Fitting

I fit an asymptotic exponential model from given problem with training data set.

```
control0 <- nls.control(maxiter = 5000, tol = 1e-05, minFactor = 1/10)
full.mod <- nls(bone ~ (beta1 - beta2 * exp(-beta3 * age)), data=train,
               start=list(beta1 = 115, beta2 = 120, beta3 = 0.9), trace=T,
               control=control0)
```

```
## 28453.36 : 115.0 120.0 0.9
## 15816.13 : 112.0642268 114.8791089 0.4614787
## 4221.839 : 109.3227563 108.8996564 0.1746354
## 3741.965 : 110.6666282 116.1033177 0.1457199
## 3725.994 : 110.7395722 116.6675732 0.1509562
## 3725.985 : 110.7602931 116.6630297 0.1508762
## 3725.985 : 110.7601249 116.6634934 0.1508789
```

```
summary(full.mod)
```

```
##
## Formula: bone ~ (beta1 - beta2 * exp(-beta3 * age))
##
## Parameters:
##      Estimate Std. Error t value Pr(>|t|)
```

```
## beta1 110.76012      2.50589  44.200 < 2e-16 ***
## beta2 116.66349      7.96750  14.642 9.78e-16 ***
## beta3   0.15088      0.02114   7.138 4.22e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 10.79 on 32 degrees of freedom
##
## Number of iterations to convergence: 6
## Achieved convergence tolerance: 7.28e-07
```

Remarks: The p-values of each case (β_1 , β_2 , and β_3) are smaller than the significance level (0.05), so the estimates of β_1 , β_2 , and β_3 are highly significant.

3.2 Problem-3(b): Hypothesis Testing

```
red.mod <- nls(bone ~ (beta1 - beta1 * exp(-beta3 * age)), data=train,
              start=list(beta1 = 140, beta3 = .1), trace=T,
              control=control0)
```

```
## 18604.11 : 140.0 0.1
## 4107.508 : 109.2341511 0.1297439
## 3791.072 : 110.9584509 0.1418749
## 3790.713 : 111.0949813 0.1415447
## 3790.713 : 111.0934313 0.1415632
## 3790.713 : 111.0935209 0.1415622
```

The aim of model fitting is to explain as much of the variation in the dependent variable as possible from information contained in the independent variables. The contributions of the independent variables to the model are measured by partitions of the total sum of squares of the difference of data and model (“analysis of variance”, ANOVA).

```
anova(full.mod, red.mod)
```

```
## Analysis of Variance Table
##
## Model 1: bone ~ (beta1 - beta2 * exp(-beta3 * age))
## Model 2: bone ~ (beta1 - beta1 * exp(-beta3 * age))
##   Res.Df Res.Sum Sq Df Sum Sq F value Pr(>F)
## 1      32      3726.0
## 2      33      3790.7 -1 -64.729  0.5559 0.4614
```

Remarks: As expected, the residual sums of squares of reduced model is bigger than the full model. By definition, the null hypothesis of the test says that the reduced model is correct. On the other hand, the alternative hypothesis says that the reduced model is too simple and that the more complex full model is more appropriate. From ANOVA table, the p-value > 0.05 indicates that:


1. We accept the null hypothesis and conclude that the reduced model is better;

2. There is no lack of fit. The reason is that the pure error sd. $\sqrt{(3790.7/33)} = 10.71$ is less than the regression standard error of 10.79.

3.3 Problem-3(c) Fitted curve

I add a fitted curve of reduced model to the scatterplot.

```
plot(dat$age, dat$bone, main="Bone vs. Age", xlab="Ages ", ylab="Bones", col="orange", pch=1)
lines(sort(train$age), fitted(red.mod)[order(train$age)], lty = 2, col = "red", lwd = 2)
```



3.4 Problem-3(d) MSE

Now I apply the reduced model to the test dataset and compute the prediction mean square error (MSE).

```
pred <- predict(red.mod, test)
pmse <- function(y, yhat) mean((y-yhat)^2)
(a <- pmse(test$bone, pred))
```

```
## [1] 301.7998
```

4 Problem-4: Local Regression Methods


4.1 Problem 4(a): KNN regression

Given the data set, I want to find a regression function such that it is the best match of our data. For the KNN regression, the algorithm uses ‘feature similarity’ to predict values of any new data points. This means that the new point is assigned a value based on how closely it resembles the points in the training set. I considered the number of neighbours as $k=1$, since it gives minimum SSE in the cross validation.

```
library("FNN")
set.seed(123)
# Fina a optimal K via 10-fold CV
SSEP <- function(yobs, yhat) sum((yobs-yhat)^2)
K <- 1:10
V <- 3
id.fold <- sample(1:V, size = NROW(train), replace=T)
SSE <- rep(0, length(K))
for(k in 1:length(K)){
  for(v in 1:V){
    train1<- train[id.fold!=v, ];
    train2<- train[id.fold==v, ];
    yhat2 <- knn.reg(train=train1, y=train1$bone, test=train2, k=K[k], algorithm="kd_tree")$pred
    SSE[k] <- (SSE[k] + SSEP(train2$bone, yhat2))
  }
}
cbind(K, SSE)
```

```
##      K      SSE
## [1,] 1    776.9259
## [2,] 2   1107.7884
## [3,] 3   1576.6320
## [4,] 4   3230.0434
## [5,] 5   5251.5361
## [6,] 6   7535.4341
## [7,] 7   9258.5638
## [8,] 8  11310.4174
## [9,] 9  13415.5427
## [10,] 10 14768.3572
```

```
fit.knn <- knn.reg(train=train, y=train$bone, k=1, algorithm=("kd_tree"))
plot(dat$age, dat$bone, main="Bone vs. Age", xlab="Ages ", ylab="Bones", col="orange", pch=1)
lines(train$age, fit.knn$pred, col="blue", lwd=2, lty = 2)
```



I then applied the fitted model to test data and obtain the prediction MSE.

```
pred2 = knn.reg(train = train, test = test, y=train$bone, k = 1)
(b = pmse(test$bone, pred2$pred))
```

```
## [1] 72.64694
```

4.2 Problem 4(b): Kernel Regression

The kernel Regression is a non parametric technique to fit our data. It does not assume any underlying distribution to estimate the regression function. That is why, kernel regression is categorized as non parametric technique. To improve the smoothness in KNN, kernel regression assigns the weights through a smooth kernel function, that is a pdf symmetric to 0.

The kernel bandwidth works as a smoothing parameter. All kernels are scaled, so the upper and lower quartiles of the kernel (viewed as a probability density) are ± 0.25 . Larger values of bandwidth make smoother estimates; smaller values of bandwidth make less smooth estimates.

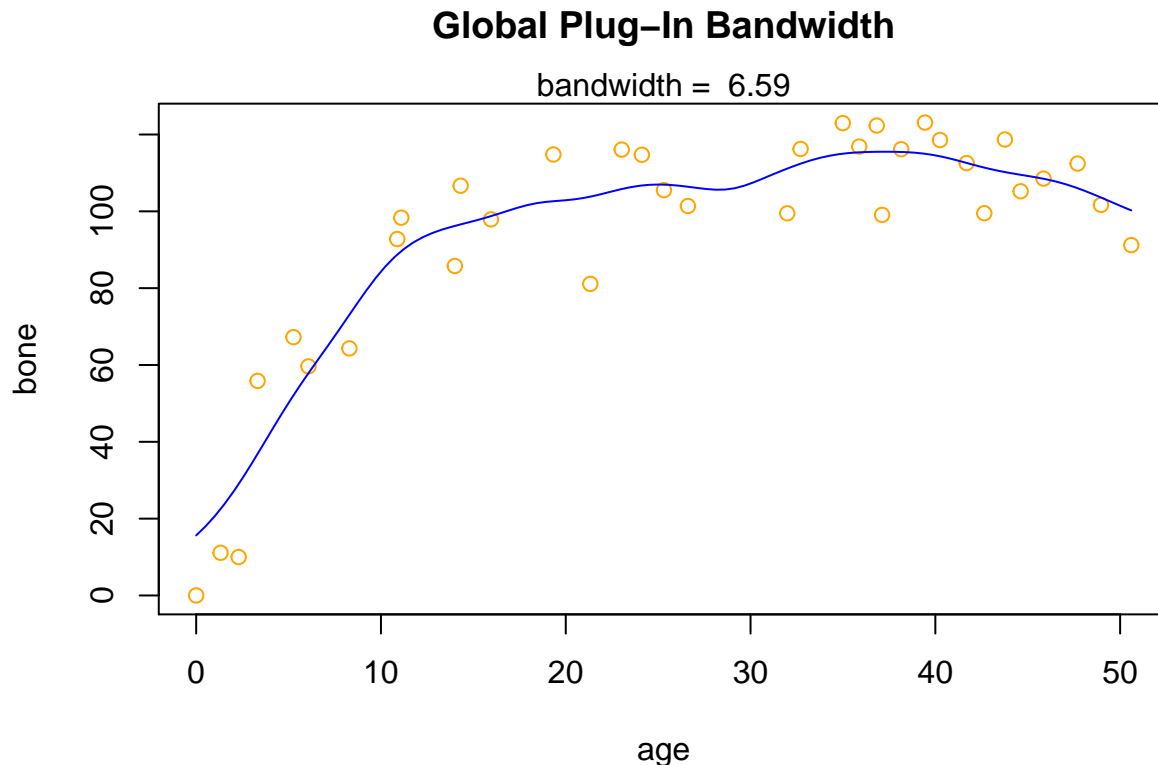
But too small bandwidths will lead to a wiggly curve, and too large ones will smooth away important details. The function `gkerns` calculates an estimator of the regression function or derivatives of the regression function with an automatically chosen global plugin bandwidth (6.59).

I use the optimal bandwidth in `ksmooth`, for which `predict()` works.

```
library(lokern)
with(train, {
  par(mfrow=c(1,1))
  plot(bone ~ age, data = train, main = "Global Plug-In Bandwidth", col="orange")
```



```
fit <- glkerns(train$age, train$bone)
lines(ksmooth(train$age, train$bone, "normal", bandwidth = fit$bandwidth), col = "blue")
mtext(paste("bandwidth = ", format(fit$bandwidth, dig = 4)))
})
```



Now I use the test dataset to predict by kernel regression technique with optimal bandwidth, and compute the prediction mean square error (MSE).

```
fit <- glkerns(train$age, train$bone)
pred3 <- predict(fit, newdata=test)

## using first column of data.frame as 'x'
(c <- pmse(test$bone, pred3$y))

## [1] 317.5978
```

4.3 Problem 4(c): Local (Cubic) Polynomial Regression

The bias problems in kernel regression can be alleviated by local polynomial regression. Here, we use EpaK kernels in the argument list. To predict the data, we use test data into the argument of xeval which is the vector of evaluation points. From lpFit we find the fitted value of bones and we estimate the MSE.

```
library(locpol)
r <- locpol(bone~age, data=train, deg=3, kernel=EpaK, bw=5)
pred4 <- locpol(bone~age, data=train, xeval=test$age, deg=3, kernel=EpaK, bw=5)
```

```
(d <-pmse(test$bone, pred4$lpFit$bone))
```

```
## [1] 870.8387
```

5 Problem 5: Regression/ Smoothing Splines

5.1 Problem 5(a): Regression Splines (Natural Cubic)

Splines fit the data very smoothly in most of the cases where polynomials would become wiggly and overfit the training data due to high variance at high degrees of polynomial.

Here, we use a Regression spline (non linear regression technique), where instead of building one model for the entire dataset, we divide the dataset into multiple bins and fit each bin with a separate model. The points where the division occurs are called Knots. The functions which we can use for modelling each piece/bin are known as Piecewise functions. There are various piecewise functions that we can use to fit these individual bins.

One problem with regression splines is that the estimates tend to display erratic behavior, i.e., they have high variance, at the boundaries of the domain of $x_1 \cdots x_n$. Even this gets worse as the order k gets larger. We can ignore this issue using natural splines, i.e., to force the piecewise polynomial function to have a lower degree to the left of the leftmost knot, and to the right of the rightmost knot. The cubic splines have continuous 1st Derivative and continuous 2nd derivative. The locations of the knots are typically quantiles of X , and the number of knots, K , is chosen by cross validation.

In the following code, we use *ns* function which generates a basis matrix for representing the family of piecewise-cubic splines with the specified sequence of interior knots, and the natural boundary conditions. These enforce the constraint that the function is linear beyond the boundary knots, which can either be supplied or default to the extremes of the data.

We then use *bs* function, which generates a basis matrix for representing the family of piecewise polynomials with the specified interior knots and degree, evaluated at the values of age variable.

After that, we compute the MSE, by predicting y using test data.

```
library(splines)
```

```
ns(data$age, df = 5)
```

```
##           1           2           3           4           5
## [1,] 0.000000e+00 0.000000e+00 0.000000000 0.00000000 0.000000000
## [2,] 2.598412e-02 0.000000e+00 -0.129418065 0.37483551 -0.245417446
## [3,] 4.473601e-04 0.000000e+00 -0.036840523 0.10670177 -0.069861244
## [4,] 1.521964e-02 4.657894e-01 0.455169213 0.12022528 -0.056403553
## [5,] 8.469135e-04 0.000000e+00 -0.045413853 0.13153283 -0.086118980
## [6,] 2.359275e-03 0.000000e+00 -0.063261915 0.18322645 -0.119964531
## [7,] 7.131540e-03 0.000000e+00 -0.089436292 0.25903569 -0.169599402
## [8,] 1.286183e-02 0.000000e+00 -0.106579388 0.30868750 -0.202108114
## [9,] 2.840420e-02 0.000000e+00 -0.132422423 0.38353708 -0.251114657
## [10,] 4.361012e-02 0.000000e+00 -0.146841120 0.42529817 -0.278457050
## [11,] 7.370391e-02 0.000000e+00 -0.163018267 0.47215229 -0.309134021
```

```

## [12,] 1.108555e-01 0.000000e+00 -0.172343138 0.49916005 -0.326816915
## [13,] 1.437199e-01 3.768559e-06 -0.175444957 0.50814390 -0.332698942
## [14,] 2.443530e-01 1.368727e-03 -0.170997564 0.49526285 -0.324265283
## [15,] 2.575969e-01 1.830388e-03 -0.169488948 0.49089342 -0.321404472
## [16,] 3.200564e-01 5.259625e-03 -0.160413109 0.46460693 -0.304193820
## [17,] 4.412133e-01 2.029041e-02 -0.135135305 0.39139444 -0.256259135
## [18,] 4.603060e-01 2.413151e-02 -0.130271438 0.37730715 -0.247035710
## [19,] 5.514655e-01 5.241463e-02 -0.102943401 0.29815654 -0.195213139
## [20,] 5.574020e-01 5.511034e-02 -0.100858505 0.29211802 -0.191259517
## [21,] 6.001247e-01 8.047545e-02 -0.083994680 0.24327507 -0.159280390
## [22,] 6.433281e-01 1.350302e-01 -0.058734745 0.17011433 -0.111379590
## [23,] 6.505857e-01 1.658442e-01 -0.048675998 0.14098103 -0.092305035
## [24,] 6.427136e-01 2.292052e-01 -0.033796453 0.09829081 -0.064354308
## [25,] 6.249461e-01 2.753867e-01 -0.025449008 0.07609580 -0.049822486
## [26,] 5.927989e-01 3.338996e-01 -0.015911296 0.05467884 -0.035800082
## [27,] 5.591125e-01 3.830212e-01 -0.007878587 0.04101751 -0.026855546
## [28,] 5.020618e-01 4.525417e-01 0.004959264 0.02706912 -0.017723068
## [29,] 4.310047e-01 5.243460e-01 0.022415958 0.01857956 -0.012164662
## [30,] 3.490284e-01 5.917798e-01 0.047447529 0.01649881 -0.010802327
## [31,] 2.884283e-01 6.312918e-01 0.071412075 0.01913850 -0.012530623
## [32,] 2.544397e-01 6.490497e-01 0.087827237 0.02210095 -0.014470238
## [33,] 1.649585e-01 6.744128e-01 0.148253390 0.03583147 -0.023460072
## [34,] 1.472617e-01 6.738477e-01 0.165113366 0.03990276 -0.026125615
## [35,] 7.732498e-02 6.347553e-01 0.264505736 0.06445437 -0.041040412
## [36,] 5.775462e-02 6.060475e-01 0.307102386 0.07559033 -0.046494881
## [37,] 3.616974e-02 5.556611e-01 0.368252580 0.09263117 -0.052714601
## [38,] 1.785032e-02 4.813338e-01 0.441638692 0.11556437 -0.056387168
## [39,] 9.615748e-03 4.243705e-01 0.488579238 0.13258263 -0.055148081
## [40,] 4.154474e-03 3.619771e-01 0.531895654 0.15132987 -0.049357118
## [41,] 3.046245e-03 3.430080e-01 0.543351978 0.15713138 -0.046537613
## [42,] 6.870692e-04 2.765064e-01 0.576371365 0.17820246 -0.031767332
## [43,] 5.634986e-06 2.012557e-01 0.596022456 0.20457166 -0.001855406
## [44,] 0.000000e+00 1.604666e-01 0.593814215 0.22112603 0.024593205
## [45,] 0.000000e+00 1.023213e-01 0.561624601 0.25015263 0.085901477
## [46,] 0.000000e+00 7.281170e-02 0.522448779 0.26915543 0.135584094
## [47,] 0.000000e+00 4.627512e-02 0.461025230 0.29114264 0.201557010
## [48,] 0.000000e+00 3.106049e-02 0.404202459 0.30778962 0.256947427
## [49,] 0.000000e+00 1.548910e-02 0.308398306 0.33186805 0.344244544
## [50,] 0.000000e+00 1.397554e-02 0.295048831 0.33496354 0.356012097
## [51,] 0.000000e+00 3.541967e-03 0.143592760 0.36741705 0.485448218
## [52,] 0.000000e+00 6.376943e-04 0.020200528 0.39164295 0.587518830
## [53,] 0.000000e+00 2.994617e-04 -0.016437144 0.39862873 0.617508947
## [54,] 0.000000e+00 0.000000e+00 -0.146042808 0.42298601 0.723056801
## attr(,"degree")
## [1] 3
## attr(,"knots")
##      20%      40%      60%      80%
## 8.740936 19.573965 30.048844 39.764521

```


```
## attr("Boundary.knots")
## [1] 0.0000 50.6041
## attr("intercept")
## [1] FALSE
## attr("class")
## [1] "ns"      "basis"    "matrix"

fm1 <- lm(bone ~ bs(age, df = 5), data = train)
summary(fm1)

##
## Call:
## lm(formula = bone ~ bs(age, df = 5), data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -23.081  -5.719   1.808   8.228  16.291
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -5.449     7.660  -0.711   0.483
## bs(age, df = 5)1  92.762    15.909   5.831 2.54e-06 ***
## bs(age, df = 5)2 107.881    12.434   8.676 1.49e-09 ***
## bs(age, df = 5)3 121.023    15.094   8.018 7.66e-09 ***
## bs(age, df = 5)4 120.777    11.478  10.522 2.05e-11 ***
## bs(age, df = 5)5  99.357    11.737   8.465 2.50e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 10.41 on 29 degrees of freedom
## Multiple R-squared:  0.9111, Adjusted R-squared:  0.8958
## F-statistic: 59.44 on 5 and 29 DF,  p-value: 2.371e-14

par(mfrow=c(1,1))
plot(bone~age, data=train, xlab = "ages", ylab = "bones", main="Deer Jaws: Natural Cubic Spline")
#spd <- seq(min(train$age), max(train$age), len = 35)
lines(sort(train$age), fm1$fitted.values[order(train$age)], lty = 2, col = "red", lwd = 2)
```

Deer Jaws: Natural Cubic Splines



```
pred5 <- predict(fm1, test)
(e <- pmse(test$bone, pred5))
```

```
## [1] 357.6741
```

5.2 Problem 5(b): Smoothing Splines

Smoothing splines are a regularized regression over the natural spline basis. In smoothing splines, we have a Knot at every unique value of x_1, \dots, x_n . It circumvents the problem of knot selection (as they just use the inputs as knots), and simultaneously, they control for overfitting by shrinking the coefficients of the estimated function (in its basis expansion). So it does not require the selection of the number of Knots, but require selection of only a Roughness Penalty (λ) which accounts for the wiggleness (fluctuations) and controls the roughness of the function and variance of the Model.

The smaller the λ , the more wiggly and fluctuating the function is. As λ , approaches ∞ , the function $g(x)$ becomes linear.


```
library(splines)
plot(train$age, train$bone, main = "Modeling Deer Jaws Data via Smoothing Splines", col="orange",
     (jaws.spl <- smooth.spline(data$age, data$bone, cv=TRUE)))
```

```
## Call:
## smooth.spline(x = data$age, y = data$bone, cv = TRUE)
##
## Smoothing Parameter spar= 0.8004763 lambda= 0.002580874 (11 iterations)
## Equivalent Degrees of Freedom (Df): 5.285307
## Penalized Criterion (RSS): 8868.881
```

```
## PRESS(1.o.o. CV): 489.4116
```

```
lines(jaws.spl, col = "blue")
```

Modeling Deer Jaws Data via Smoothing Splines



```
pred6 <- predict(jaws.spl, test)
(f <- pmse(test$bone, pred6$y$bone))
```

```
## [1] 3215.662
```

I determine the tuning parameter as $\lambda = 0.002580874$ from the cross-validation. That is why, I use $CV = TRUE$ in the code.

6 Problem-6

```
(tabulate <- data.frame(
  Methods = c("NLS", "KNN", "Kernel Regression", "local (cubic) Poly.", "Regression Spline", "Smooth Spline"),
  MSE = c(a, b, c, d, e, f)
))
```

```
##           Methods      MSE
## 1           NLS 301.79976
## 2           KNN  72.64694
## 3  Kernel Regression 317.59778
## 4 local (cubic) Poly. 870.83866
## 5  Regression Spline 357.67407
## 6   Smooth spline 3215.66187
```

The KNN mehtod gives minimum MSE and thus it gives favorable results.