

Project-VI

Md Al Masum Bhuiyan

December 12, 2018

Contents

1 Problem-1: Data Preparation	2
2 Problem-2: Exploratory Data Analysis	4
2.1 Frequency Distribution of the Target variable Outcome	6
2.2 Density plot	7
2.3 Correlation Matrix	7
3 Problem-3: Data Partition	10
4 Problem-4: Logistic Regression	10
5 Problem-5: Random Forest	15
6 Problem-3: Artificial Neural Network	18
6.1 ANN1: Single hidden layer	18
6.2 ANN2: Two hidden layers	22
6.3 ANN3: Three hidden layers	25
7 Support Vector Machine	29
7.1 SVM-Liner	29
7.2 SVM-Linear-grid	32
7.3 SVM-Radial	34
7.4 SVM-Radial-grid	37
7.5 SVM: Radial Basis kernel "Gaussian"	42
8 Comparison among the fitted models	45

1 Problem-1: Data Preparation

I begin the project VI by bringing the Climate Model Simulation Crashes data into Rmd. The data comes from the UCI machine learning repository. It contains 540 observations and 21 variables.

```
dat <- read.table(file=
  "http://archive.ics.uci.edu/ml/machine-learning-databases/00252/pop_failures.dat",
  header=TRUE)
dim(dat); head(dat); anyNA(dat);
```

```
## [1] 540 21

##   Study Run vconst_corr vconst_2 vconst_3 vconst_4 vconst_5 vconst_7
## 1     1   1  0.85903621 0.9278245 0.25286562 0.2988383 0.1705213 0.7359360
## 2     1   2  0.60604103 0.4577284 0.35944842 0.3069574 0.8433308 0.9348507
## 3     1   3  0.99759978 0.3732385 0.51739936 0.5049925 0.6189033 0.6055708
## 4     1   4  0.78340786 0.1040553 0.19753270 0.4218372 0.7420557 0.4908279
## 5     1   5  0.40624953 0.5131993 0.06181158 0.6358367 0.8447975 0.4415022
## 6     1   6  0.04137947 0.6290259 0.30338011 0.8134076 0.2228171 0.9712060
##      ah_corr ah_bolus slm_corr efficiency_factor tidal_mix_max
## 1 0.428325428 0.5679469 0.47436960      0.2456749      0.10422587
## 2 0.444572488 0.8280149 0.29661775      0.6168699      0.97578558
## 3 0.746225330 0.1959283 0.81566694      0.6793550      0.80341308
## 4 0.005525437 0.3921233 0.01001489      0.4714627      0.59787890
## 5 0.191926451 0.4875461 0.35853358      0.5515432      0.74387652
## 6 0.609778290 0.6478039 0.73791387      0.4409432      0.03598237
##   vertical_decay_scale convect_corr bckgrnd_vdc1 bckgrnd_vdc_ban
## 1           0.8690907    0.99751850    0.4486201    0.30752179
## 2           0.9143437    0.84524714    0.8641519    0.34671269
## 3           0.6439952    0.71844113    0.9247751    0.31537141
## 4           0.7616588    0.36275056    0.9128191    0.97797118
## 5           0.3123494    0.65022283    0.5222610    0.04354476
## 6           0.6158677    0.01748718    0.9323195    0.32931804
##   bckgrnd_vdc_eq bckgrnd_vdc_psim Prandtl outcome
## 1           0.8583104           0.7969972 0.8698930      0
## 2           0.3565734           0.4384472 0.5122561      1
## 3           0.2506424           0.2856355 0.3658580      1
## 4           0.8459212           0.6994309 0.4759867      1
## 5           0.3766601           0.2800978 0.1322829      1
## 6           0.9541228           0.1353789 0.2948048      1

## [1] FALSE
```

This dataset contains records of simulation crashes encountered during climate model uncertainty quantification (UQ) ensembles. I remove the first two columns and do the exploratory data analysis.

I remove the first two columns.


```
data <- dat[, c(-1,-2)]
```

Now I check the missing values.


```
##### Missing values by visualization #####
```

```
library(naniar)
```

```
vis_miss(data)
```



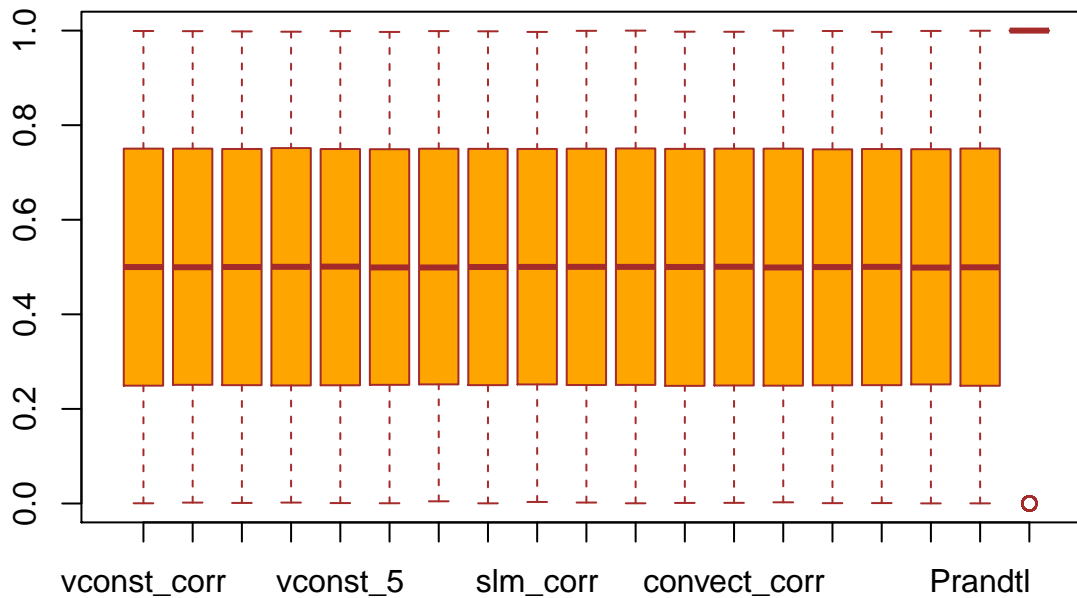
```
gg_miss_var(data)
```



2 Problem-2: Exploratory Data Analysis


```
boxplot(data, main = "Boxplot", horizontal = F, col="orange", border="brown")
```

Boxplot



Remarks: The box plot shows that the variability of each predictors are same. Now I present another boxplot for outcome resonse and a predictor vconst_7.

```
library(ggplot2)
ggplot(aes(x=outcome, fill = vconst_7), data = data)+geom_bar()
```




2.1 Frequency Distribution of the Target variable Outcome

The following code shows the frequency, percentage, cumulative percentage of the target variable Outcome..

```
library("funModeling")

## Loading required package: Hmisc
## Loading required package: lattice
## Loading required package: survival
## Loading required package: Formula
##
## Attaching package: 'Hmisc'
## The following objects are masked from 'package:base':
##
##   format.pval, units
## funModeling v.1.6.8 :)
## Examples and tutorials at livebook.datascienceheroes.com

freq(data$outcome)
```




```
##   var frequency percentage cumulative_perc
```

```
## 1    1      494      91.48      91.48
## 2    0       46       8.52     100.00
```

2.2 Density plot

```
library(ggplot2)
ggplot(aes(vconst_corr), data=data) + geom_density()
```



2.3 Correlation Matrix

From the correlation matrix, we see the correlation among the variables. For example, `Vconst_1` and `Vconst_2` have the correlation values 0.004039037, that is, positively correlated.

```
cor(data, method = "pearson", use = "complete.obs")
```

```
##          vconst_corr    vconst_2    vconst_3
## vconst_corr    1.000000000  0.0040390372  0.0093312390
## vconst_2       0.004039037  1.0000000000 -0.0004564662
## vconst_3       0.009331239 -0.0004564662  1.0000000000
## vconst_4      -0.018294219 -0.0006144773  0.0098992508
## vconst_5       0.018880130 -0.0082917043  0.0062887724
## vconst_7       0.001543507 -0.0243790678 -0.0015865605
## ah_corr        0.003713983 -0.0051821138  0.0199408652
```

```

## ah_bolus          -0.012735004  0.0041791244  0.0044018146
## slm_corr          0.002336425 -0.0138597114 -0.0076953527
## efficiency_factor  0.010617309 -0.0110718393  0.0071000071
## tidal_mix_max     -0.014205434  0.0197060654 -0.0094284119
## vertical_decay_scale -0.008991676  0.0016229327 -0.0247021845
## convect_corr      -0.002980494  0.0026084619 -0.0206373211
## bckgrnd_vdc1      -0.002132558 -0.0147164701 -0.0042641402
## bckgrnd_vdc_ban   -0.002098679  0.0043858772 -0.0052104219
## bckgrnd_vdc_eq     0.015972793  0.0059985634 -0.0005588026
## bckgrnd_vdc_psim  -0.016631072  0.0042022103  0.0047712303
## Prandtl           -0.001466642  0.0091410713 -0.0013344059
## outcome           -0.304786983 -0.3023878482  0.0002271854
##                  vconst_4    vconst_5    vconst_7    ah_corr
## vconst_corr       -0.0182942192  0.018880130  1.543507e-03  0.003713983
## vconst_2          -0.0006144773 -0.008291704 -2.437907e-02 -0.005182114
## vconst_3           0.0098992508  0.006288772 -1.586561e-03  0.019940865
## vconst_4           1.0000000000  0.020504177  2.193088e-02  0.001804725
## vconst_5           0.0205041766  1.0000000000  5.887138e-03 -0.003046706
## vconst_7           0.0219308802  0.005887138  1.000000e+00 -0.016770369
## ah_corr            0.0018047254 -0.003046706 -1.677037e-02  1.000000000
## ah_bolus          -0.0023344776  0.012453140 -2.164373e-02 -0.035497640
## slm_corr           -0.0017310089  0.003633894  1.243988e-03 -0.005119394
## efficiency_factor  -0.0047528309  0.001076585  1.512107e-02  0.009603721
## tidal_mix_max      0.0183200928  0.021353663  7.453114e-05 -0.006831726
## vertical_decay_scale -0.0100038424 -0.016311669  1.528827e-02  0.016502884
## convect_corr       -0.0067617001  0.021379515  7.035600e-03  0.002921182
## bckgrnd_vdc1       0.0204418668  0.009893622 -3.640898e-03  0.012447485
## bckgrnd_vdc_ban    -0.0010800378 -0.019179395 -7.897314e-03 -0.003368375
## bckgrnd_vdc_eq     -0.0092617101 -0.020752192 -6.575539e-03  0.007050740
## bckgrnd_vdc_psim   -0.0171472800 -0.009324311  1.320325e-02  0.002442509
## Prandtl            0.0050528658  0.012265166  8.411513e-03 -0.002380732
## outcome            0.0722970410  0.054390262  4.864631e-02  0.017048998
##                  ah_bolus    slm_corr efficiency_factor
## vconst_corr       -0.0127350037  0.002336425    0.010617309
## vconst_2           0.0041791244 -0.013859711    -0.011071839
## vconst_3           0.0044018146 -0.007695353    0.007100007
## vconst_4          -0.0023344776 -0.001731009    -0.004752831
## vconst_5           0.0124531397  0.003633894    0.001076585
## vconst_7          -0.0216437289  0.001243988    0.015121070
## ah_corr            -0.0354976401 -0.005119394    0.009603721
## ah_bolus           1.0000000000 -0.009403074    0.012259861
## slm_corr           -0.0094030741  1.000000000    0.008759613
## efficiency_factor  0.0122598609  0.008759613    1.000000000
## tidal_mix_max      0.0120047990  0.002574915    -0.017925605
## vertical_decay_scale -0.0039467690  0.002272225    0.018009324
## convect_corr       -0.0193073756  0.002633296    0.011924616
## bckgrnd_vdc1       -0.0106420189 -0.003043485    -0.034025929
## bckgrnd_vdc_ban    0.0048663286  0.006022810    0.003392878

```



```

## bckgrnd_vdc_eq      0.0323978926 -0.008446854      0.009924829
## bckgrnd_vdc_psim    0.0002592278 -0.002300814      -0.005241170
## Prandtl             0.0070551148  0.014280916      -0.004465106
## outcome             0.0038954358  0.048863670      -0.032364442
##                    tidal_mix_max vertical_decay_scale convect_corr
## vconst_corr        -1.420543e-02      -0.008991676 -0.002980494
## vconst_2            1.970607e-02      0.001622933  0.002608462
## vconst_3           -9.428412e-03      -0.024702184 -0.020637321
## vconst_4            1.832009e-02      -0.010003842 -0.006761700
## vconst_5            2.135366e-02      -0.016311669  0.021379515
## vconst_7            7.453114e-05      0.015288275  0.007035600
## ah_corr            -6.831726e-03      0.016502884  0.002921182
## ah_bolus            1.200480e-02      -0.003946769 -0.019307376
## slm_corr            2.574915e-03      0.002272225  0.002633296
## efficiency_factor   -1.792561e-02      0.018009324  0.011924616
## tidal_mix_max       1.000000e+00      -0.004328110  0.002345084
## vertical_decay_scale -4.328110e-03      1.000000000  0.010967371
## convect_corr        2.345084e-03      0.010967371  1.000000000
## bckgrnd_vdc1        -1.686416e-03      -0.008099906 -0.007046943
## bckgrnd_vdc_ban     1.761794e-02      -0.001098075 -0.013650361
## bckgrnd_vdc_eq     -9.566276e-03      0.007049969 -0.016532850
## bckgrnd_vdc_psim    -2.260626e-03      -0.015983717  0.006300329
## Prandtl            -1.191322e-02      0.006292207 -0.005097606
## outcome            -1.507052e-02      -0.045862217 -0.192892806
##                    bckgrnd_vdc1 bckgrnd_vdc_ban bckgrnd_vdc_eq
## vconst_corr        -0.0021325579 -0.0020986795  0.0159727928
## vconst_2           -0.0147164701  0.0043858772  0.0059985634
## vconst_3           -0.0042641402 -0.0052104219 -0.0005588026
## vconst_4            0.0204418668 -0.0010800378 -0.0092617101
## vconst_5            0.0098936216 -0.0191793953 -0.0207521918
## vconst_7           -0.0036408981 -0.0078973136 -0.0065755386
## ah_corr            0.0124474852 -0.0033683750  0.0070507404
## ah_bolus           -0.0106420189  0.0048663286  0.0323978926
## slm_corr           -0.0030434855  0.0060228097 -0.0084468541
## efficiency_factor   -0.0340259293  0.0033928784  0.0099248294
## tidal_mix_max      -0.0016864159  0.0176179416 -0.0095662761
## vertical_decay_scale -0.0080999062 -0.0010980751  0.0070499692
## convect_corr        -0.0070469427 -0.0136503606 -0.0165328497
## bckgrnd_vdc1        1.0000000000  0.0003010617  0.0100984372
## bckgrnd_vdc_ban     0.0003010617  1.0000000000 -0.0019247142
## bckgrnd_vdc_eq      0.0100984372 -0.0019247142  1.0000000000
## bckgrnd_vdc_psim    0.0058160811  0.0163438907  0.0191253312
## Prandtl            -0.0028796392  0.0037739752 -0.0008064784
## outcome             0.1842179984 -0.0283649869  0.0785044287
##                    bckgrnd_vdc_psim Prandtl outcome
## vconst_corr        -0.0166310716 -0.0014666423 -0.3047869831
## vconst_2            0.0042022103  0.0091410713 -0.3023878482
## vconst_3            0.0047712303 -0.0013344059  0.0002271854

```

## vconst_4	-0.0171472800	0.0050528658	0.0722970410
## vconst_5	-0.0093243109	0.0122651659	0.0543902621
## vconst_7	0.0132032517	0.0084115132	0.0486463050
## ah_corr	0.0024425089	-0.0023807322	0.0170489975
## ah_bolus	0.0002592278	0.0070551148	0.0038954358
## slm_corr	-0.0023008138	0.0142809159	0.0488636704
## efficiency_factor	-0.0052411699	-0.0044651060	-0.0323644420
## tidal_mix_max	-0.0022606260	-0.0119132212	-0.0150705214
## vertical_decay_scale	-0.0159837171	0.0062922071	-0.0458622169
## convect_corr	0.0063003294	-0.0050976057	-0.1928928063
## bckgrnd_vdc1	0.0058160811	-0.0028796392	0.1842179984
## bckgrnd_vdc_ban	0.0163438907	0.0037739752	-0.0283649869
## bckgrnd_vdc_eq	0.0191253312	-0.0008064784	0.0785044287
## bckgrnd_vdc_psim	1.0000000000	0.0127673739	0.0576844108
## Prandtl	0.0127673739	1.0000000000	-0.0269417345
## outcome	0.0576844108	-0.0269417345	1.0000000000

3 Problem-3: Data Partition


Now I randomly split the data D into the training set D1 and the test set D2 with a ratio of approximately 2:1 on the sample size. I used `set.seed()` to fix the random seed so that the results are easily reproducible.

```
set.seed(123)
n <- nrow(data)
split_data <- sample(x=1:2, size = n, replace=TRUE, prob=c(0.67, 0.33))
train <- data[split_data == 1, ]
test <- data[split_data == 2, ]
y.train <- train$outcome
yobs <- test$outcome
```

4 Problem-4: Logistic Regression


Here, I fit the regularized logistic regression using the training data D1 with Lasso model. In `glmnet` function, the family argument specify that we want a “binomial” model which tells `glmnet()` to fit a logistic function to the data.

```
library(glmnet)
formula0 <- outcome~.
X <- model.matrix (as.formula(formula0), data = train)
y <- train$outcome
fit.lasso <- glmnet(x=X, y=y, family="binomial", alpha=1,
                    lambda.min = 1e-4, nlambda = 300, standardize=T, thresh =
                    1e-07, maxit=3000)
plot(fit.lasso)
```



Using cross validation to determine the optimal tuning parameter.

```
CV <- cv.glmnet(x=X, y=train$outcome, family="binomial", alpha = 1,
               lambda.min = 1e-4, nlambda = 200, standardize = T, thresh = 1e-07,
               maxit=1000)
plot(CV)
```



I select the best tuning parameter (λ) and use it in the final model.

```
b.lambda <- CV$lambda.1se; b.lambda

## [1] 0.02210913

fit.best <- glmnet(x=X, y=train$outcome, family="binomial", alpha = 1,
                  lambda=b.lambda, standardize = T, thresh = 1e-07,
                  maxit=1000)

fit.best$beta

## 19 x 1 sparse Matrix of class "dgCMatrix"
##                               s0
## (Intercept)                   .
## vconst_corr                  -3.910770
## vconst_2                     -4.466996
## vconst_3                      .
## vconst_4                      .
## vconst_5                      .
## vconst_7                      .
## ah_corr                      .
## ah_bolus                     .
## slm_corr                     .
## efficiency_factor            .
## tidal_mix_max                .
## vertical_decay_scale         .
## convect_corr                 -1.600313
## bckgrnd_vdc1                 1.296857
## bckgrnd_vdc_ban              .
## bckgrnd_vdc_eq               .
## bckgrnd_vdc_psim             .
## Prandtl                      .
```

Here, I use the best lambda with fitted model to predict the test data.

```
X.test <- model.matrix(as.formula(formula0), data = test)
pred <- predict(fit.best, newx = X.test, s=b.lambda, type="response")
dim(pred)
```

```
## [1] 173  1
```

The missclassification rate is:

```
(miss.rate_a <- mean(yobs != (pred > 0.5)))
```

```
## [1] 0.07514451
```

The Mean squared error is as:

```
MSE.a <- mean((yobs-pred)^2)
MSE.a
```

```
## [1] 0.05003158
```

Plotting ROC curve of the fit.best model.

```
library(cvAUC)
AUC <- ci.cvAUC(predictions = pred, labels = yobs, folds=1:NROW(test), confidence = 0.95)
AUC


## $cvAUC
## [1] 0.9067511
##
## $se
## [1] 0.06063456
##
## $ci
## [1] 0.7879095 1.0000000
##
## $confidence
## [1] 0.95

(auc.ci <- round(AUC$ci, digits = 3))

## [1] 0.788 1.000

library(verification)
mod.glm <- verify(obs = yobs, pred = pred)

## If baseline is not included, baseline values will be calculated from the sample obs.
roc.plot(mod.glm, plot.thres=NULL)
text(x=0.7, y=0.2, paste("Area under ROC = ", round(AUC$cvAUC, digits = 3), "with 95% CI (",
                        auc.ci[1], ",", auc.ci[2], ").", sep = " "), col="blue", cex=1.2)
```



The accuracy for logistic regression is 90.07%.

```
library(caret)
```

```
##
## Attaching package: 'caret'
## The following object is masked from 'package:survival':
##
##      cluster
pred1 <- ifelse(pred>0.5, 1, 0)
confusionMatrix(factor(pred1), factor(test$outcome))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0     3     1
##           1    12   157
##
##               Accuracy : 0.9249
##               95% CI : (0.8749, 0.9594)
##       No Information Rate : 0.9133
##       P-Value [Acc > NIR] : 0.354468
##
##               Kappa : 0.2899
```

```
## McNemar's Test P-Value : 0.005546
##
##          Sensitivity : 0.20000
##          Specificity : 0.99367
##          Pos Pred Value : 0.75000
##          Neg Pred Value : 0.92899
##          Prevalence : 0.08671
##          Detection Rate : 0.01734
##          Detection Prevalence : 0.02312
##          Balanced Accuracy : 0.59684
##
##          'Positive' Class : 0
##
```

5 Problem-5: Random Forest

I fit the random forest model with training data.

```
library(randomForest)
set.seed(123)
rf_fit <- randomForest(as.factor(outcome) ~.,
                       data=train,
                       importance=TRUE,
                       proximity=TRUE,
                       ntree=200)
rf_yhat <- predict(rf_fit, newdata=test, type="prob")[, 2]
```

Variable importance plot:

I plot the importance of variables using Mean Decrease Accuracy and Mean Decrease Gini. I see that the vconst_2 and vconst_corr are most important two variables from plots.


```
round(importance(rf_fit), 2)
```

	0	1	MeanDecreaseAccuracy	MeanDecreaseGini
## vconst_corr	9.48	7.11	10.10	8.46
## vconst_2	13.29	7.28	12.18	9.27
## vconst_3	-0.95	0.99	0.60	2.03
## vconst_4	0.34	-0.58	-0.36	2.29
## vconst_5	-1.08	2.05	1.50	1.77
## vconst_7	-1.10	1.37	0.93	2.03
## ah_corr	-1.85	-1.34	-2.03	2.02
## ah_bolus	0.23	0.05	0.22	1.91
## slm_corr	-0.88	0.01	-0.33	1.80
## efficiency_factor	-1.94	-0.26	-1.02	2.10
## tidal_mix_max	-2.25	-0.27	-1.44	1.67
## vertical_decay_scale	-0.13	0.83	0.80	1.75
## convect_corr	3.69	1.59	3.64	5.74

```
## bckgrnd_vdc1      6.30  2.21      5.07      4.17
## bckgrnd_vdc_ban   0.56  0.96      1.27      2.34
## bckgrnd_vdc_eq    2.53  2.12      3.17      2.06
## bckgrnd_vdc_psim  0.86 -1.40     -0.88      2.08
## Prandtl          -0.50 -0.03     -0.18      2.47
```

```
varImpPlot(rf_fit, main="Variable Importance Ranking", col="red")
```

Variable Importance Ranking



The missclassification rate for random forest is:

```
(miss.rate_b <- mean(yobs != (rf_yhat > 0.5)))
```

```
## [1] 0.0867052
```

Mean squared error:

```
MSE.b <- mean((yobs-rf_yhat)^2)
MSE.b
```

```
## [1] 0.06026633
```

```
##### AUC #####
```

```
library(cvAUC)
```

```
rf_AUC <- ci.cvAUC(predictions = rf_yhat, labels = yobs, folds=1:NROW(test), confidence = 0.95)
rf_AUC
```

```
## $cvAUC
```

```
## [1] 0.8563291
```

```
##
```



```
## $se
## [1] 0.05479113
##
## $ci
## [1] 0.7489405 0.9637178
##
## $confidence
## [1] 0.95


(rf_auc.ci <- round(rf_AUC$ci, digits = 3))

## [1] 0.749 0.964

library(verification)
mod.rf <- verify(obs = yobs, pred = rf_yhat)

## If baseline is not included, baseline values will be calculated from the sample obs.
roc.plot(mod.rf, plot.thres=NULL)
text(x=0.7, y=0.2, paste("Area under ROC = ", round(rf_AUC$cvAUC, digits = 3), "with 95% CI ( ",
                        rf_auc.ci[1], ",", rf_auc.ci[2], ").", sep = " "), col="blue", cex = 1)
```

ROC Curve



The accuracy is 85.60% for random forest model.

```
rf_yhat1 <- ifelse(rf_yhat>0.5, 1, 0)
confusionMatrix(factor(rf_yhat1), factor(test$outcome))
```

```
## Warning in confusionMatrix.default(factor(rf_yhat1), factor(test$outcome)):
```

```
## Levels are not in the same order for reference and data. Refactoring data
## to match.

## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0    0    0
##           1   15 158
##
##           Accuracy : 0.9133
##           95% CI : (0.861, 0.9507)
##    No Information Rate : 0.9133
##    P-Value [Acc > NIR] : 0.5681598
##
##           Kappa : 0
##    Mcnemar's Test P-Value : 0.0003006
##
##           Sensitivity : 0.00000
##           Specificity : 1.00000
##    Pos Pred Value :      NaN
##    Neg Pred Value : 0.91329
##           Prevalence : 0.08671
##    Detection Rate : 0.00000
##    Detection Prevalence : 0.00000
##    Balanced Accuracy : 0.50000
##
##           'Positive' Class : 0
##
```

6 Problem-3: Artificial Neural Network

6.1 ANN1: Single hidden layer

I first prepare the data for artificial neural network.

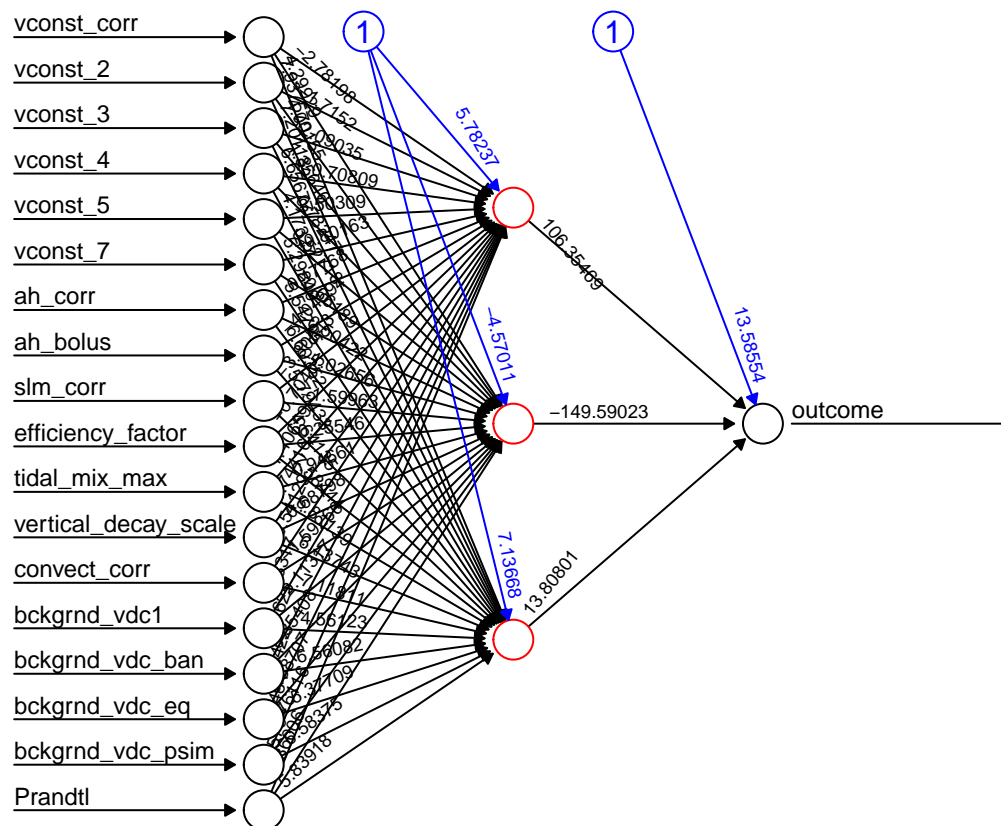
```
# DATA PREPARATION - NEED TO DEAL WITH CATEGORICAL PREDICTORS
X.dat <- as.data.frame(model.matrix(outcome~.-1, data=data))
dat1 <- data.frame(cbind(X.dat, outcome=data$outcome))

set.seed(123)
n <- nrow(dat1)
split_data <- sample(x=1:2, size = n, replace=TRUE, prob=c(0.67, 0.33))
train1 <- dat1[split_data == 1, ]
test1 <- dat1[split_data == 2, ]
yobs <- test1$outcome
test1 <- test1[, -19]
```

I use 3 neurons for single layer.

```
library(neuralnet);
options(digits=3)
net1 <- neuralnet(outcome~vconst_corr+vconst_2+vconst_3+
                  vconst_4+vconst_5+vconst_7+ah_corr+ah_bolus+slm_corr+efficiency_factor+
                  tidal_mix_max+vertical_decay_scale+convect_corr+bckgrnd_vdc1+
                  bckgrnd_vdc_ban+bckgrnd_vdc_eq+bckgrnd_vdc_psim+Prandtl,
                  data=train1,
                  hidden=3, #1 hidden layer, 3 units
                  act.fct='logistic', err.fct="ce", linear.output=F, likelihood=TRUE)

# PLOT THE MODEL
plot(net1, rep="best", show.weights=T, dimension=6.5, information=F, radius=.15,
     col.hidden="red", col.hidden.synapse="black", lwd=1, fontsize=9)
```



```
# PREDICTION
ypred_net1 <- compute(net1, covariate=test1)$net.result
```

The mean squared error is as:

```
MSE.c <- mean((yobs-ypred_net1)^2)
MSE.c
```

```
## [1] 0.06982044344
```

```

ypred_net11 <- ifelse(ypred_net1>0.5,1,0)
(miss.rate_c <- mean(yobs != ypred_net11))

## [1] 0.06936416185

##### AUC #####
library(cvAUC)
net_AUC <- ci.cvAUC(predictions = ypred_net1, labels =yobs, folds=1:NROW(test1), confidence = 0.95)
net_AUC

## $cvAUC
## [1] 0.9506329114
##
## $se
## [1] 0.02299775766
##
## $ci
## [1] 0.9055581346 0.9957076881
##
## $confidence
## [1] 0.95


(net_auc.ci <- round(net_AUC$ci, digits = 3))

## [1] 0.906 0.996

library(verification)
mod.net1 <- verify(obs = yobs, pred = ypred_net1)

## If baseline is not included, baseline values will be calculated from the sample obs.
roc.plot(mod.net1, plot.thres=NULL)
text(x=0.7, y=0.2, paste("Area under ROC = ", round(net_AUC$cvAUC, digits = 3), "with 95% CI (",
                        net_auc.ci[1], ",", net_auc.ci[2], ").", sep = " "), col="blue", cex = 1.2)

```



```
confusionMatrix(factor(ypred_net11), factor(test$outcome))
```

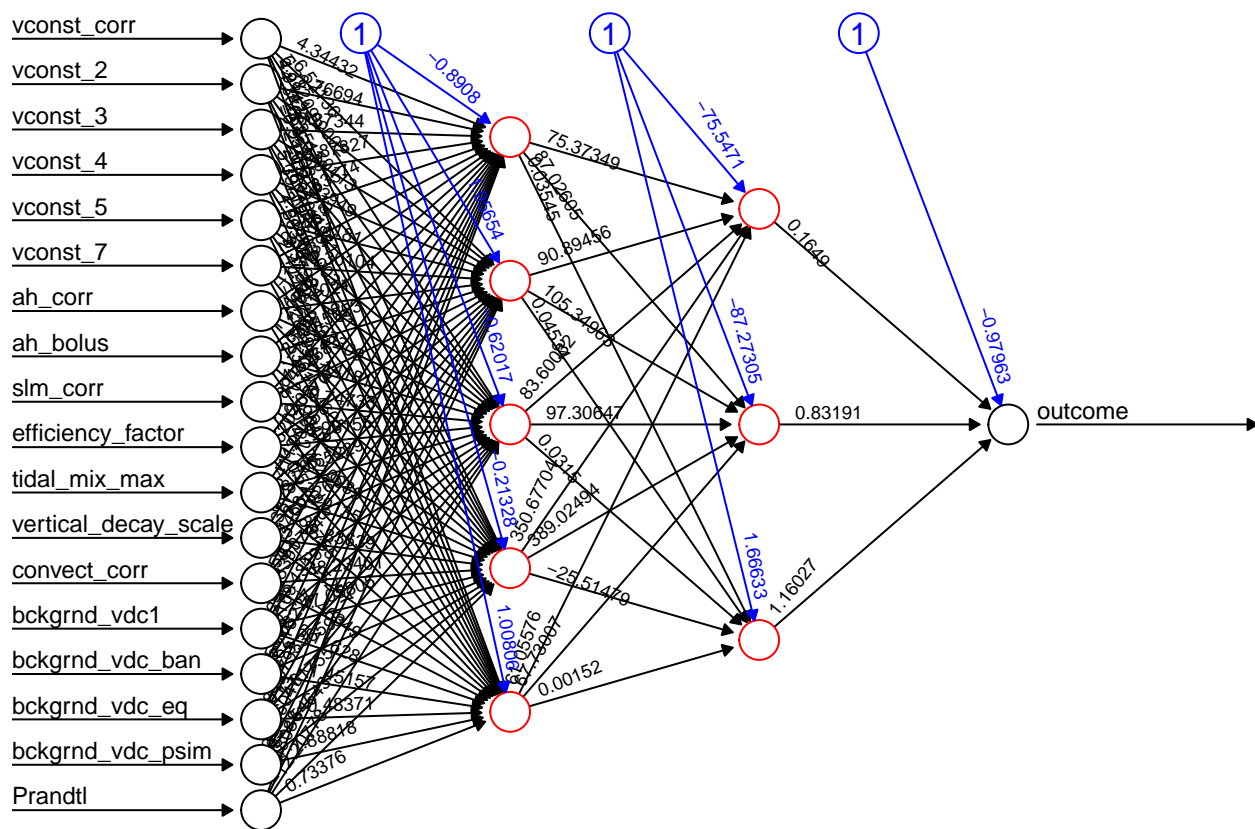
```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0    9    6
##           1    6 152
##
##           Accuracy : 0.9306358
##           95% CI   : (0.88197, 0.9636472)
##    No Information Rate : 0.9132948
##    P-Value [Acc > NIR] : 0.2563017
##
##           Kappa : 0.5620253
##    McNemar's Test P-Value : 1.0000000
##
##           Sensitivity : 0.6000000
##           Specificity : 0.9620253
##    Pos Pred Value : 0.6000000
##    Neg Pred Value : 0.9620253
##           Prevalence : 0.08670520
##    Detection Rate : 0.05202312
##    Detection Prevalence : 0.08670520
##    Balanced Accuracy : 0.78101266
```

```
##
##      'Positive' Class : 0
##
```

6.2 ANN2: Two hidden layers

I used two hidden layers(5 neurons in first layer, and 3 neurons in second layer)

```
nn <- neuralnet(outcome~vconst_corr+vconst_2+vconst_3+
                 vconst_4+vconst_5+vconst_7+ah_corr+ah_bolus+slm_corr+efficiency_factor+
                 tidal_mix_max+vertical_decay_scale+convect_corr+bckgrnd_vdc1+
                 bckgrnd_vdc_ban+bckgrnd_vdc_eq+bckgrnd_vdc_psim+Prandtl, data=train1,
                 hidden=c(5,3),linear.output=T)
plot(nn, rep="best", show.weights=T, dimension=6.5, information=F, radius=.15,
     col.hidden="red", col.hidden.synapse="black", lwd=1, fontsize=9)
```



```
# PREDICTION
ypred_nn <- compute(nn, covariate=test1)$net.result

ypred_nn1 <- ifelse(ypred_nn>0.5,1,0)
(miss.rate_d <- mean(yobs != ypred_nn1))
```

```
## [1] 0.08670520231
```


```
MSE.d <- mean((yobs-ypred_nn)^2)
MSE.d
```

```
## [1] 0.08478119453
```

Here I scaled the predicted probabilities between 0 and 1.

```
ypred_nn_scale <- scale(ypred_nn, center = min(ypred_nn), scale=max(ypred_nn)-min(ypred_nn))
hist(ypred_nn_scale)
```

Histogram of ypred_nn_scale



```
##### AUC #####
```

```
library(cvAUC)
```

```
nn_AUC <- ci.cvAUC(predictions = ypred_nn_scale, labels =yobs, folds=1:NROW(test1), confidence
nn_AUC
```

```
## $cvAUC
```

```
## [1] 0.911814346
```

```
##
```

```
## $se
```

```
## [1] 0.04784049402
```

```
##
```

```
## $ci
```

```
## [1] 0.8180487007 1.0000000000
```

```
##
```

```
## $confidence
```

```
## [1] 0.95
```


```
(nn_auc.ci <- round(nn_AUC$ci, digits = 3))

## [1] 0.818 1.000

library(verification)
mod.nn <- verify(obs = yobs, pred = ypred_nn_scale)

## If baseline is not included, baseline values will be calculated from the sample obs.
roc.plot(mod.nn, plot.thres=NULL)
text(x=0.7, y=0.2, paste("Area under ROC = ", round(nn_AUC$cvAUC, digits = 3), "with 95% CI (",
                        nn_auc.ci[1], ",", nn_auc.ci[2], ").", sep = " "), col="blue", cex = 1)
```

ROC Curve



```
confusionMatrix(factor(ypred_nn1), factor(test$outcome))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0    8    8
##           1    7 150
##
##           Accuracy : 0.9132948
##           95% CI : (0.8610262, 0.9506581)
##           No Information Rate : 0.9132948
##           P-Value [Acc > NIR] : 0.5681598
```

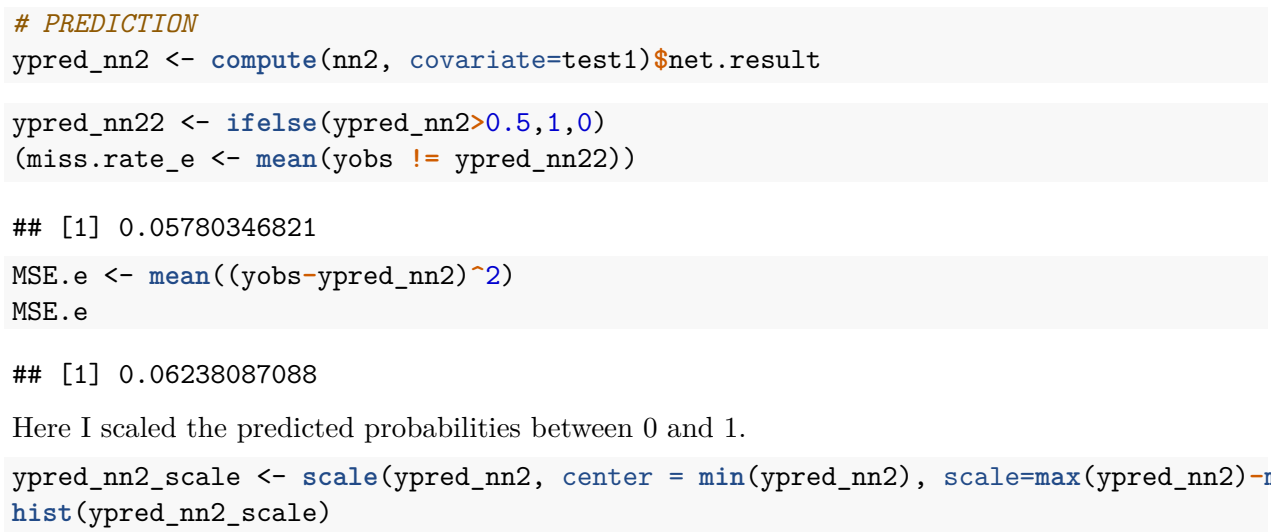


```
##
##           Kappa : 0.4685644
## McNemar's Test P-Value : 1.0000000
##
##           Sensitivity : 0.5333333
##           Specificity : 0.94936709
##           Pos Pred Value : 0.5000000
##           Neg Pred Value : 0.95541401
##           Prevalence : 0.08670520
##           Detection Rate : 0.04624277
##           Detection Prevalence : 0.09248555
##           Balanced Accuracy : 0.74135021
##
##           'Positive' Class : 0
##
```


6.3 ANN3: Three hidden layers

Here I used three hidden layers (5 neurons in first layer, 4 neurons in second layer, 2 neurons in third layer).

```
nn2 <- neuralnet(outcome~vconst_corr+vconst_2+vconst_3+
                 vconst_4+vconst_5+vconst_7+ah_corr+ah_bolus+slm_corr+efficiency_factor+
                 tidal_mix_max+vertical_decay_scale+convect_corr+bckgrnd_vdc1+
                 bckgrnd_vdc_ban+bckgrnd_vdc_eq+bckgrnd_vdc_psim+Prandtl, data=train1,
                 hidden=c(5,4,2),linear.output=T)
plot(nn2, rep="best", show.weights=T, dimension=6.5, information=F, radius=.15,
     col.hidden="red", col.hidden.synapse="black", lwd=1, fontsize=9)
```



Histogram of ypred_nn2_scale



```
##### AUC #####
library(cvAUC)
nn2_AUC <- ci.cvAUC(predictions = ypred_nn2_scale, labels = yobs, folds=1:NROW(test1), confidence=0.95)
nn2_AUC
```

```
## $cvAUC
## [1] 0.7666666667
##
## $se
## [1] 0.09526080015
##
## $ci
## [1] 0.5799589292 0.9533744041
##
## $confidence
## [1] 0.95
```

```
(nn2_auc.ci <- round(nn2_AUC$ci, digits = 3))
```

```
## [1] 0.580 0.953
```

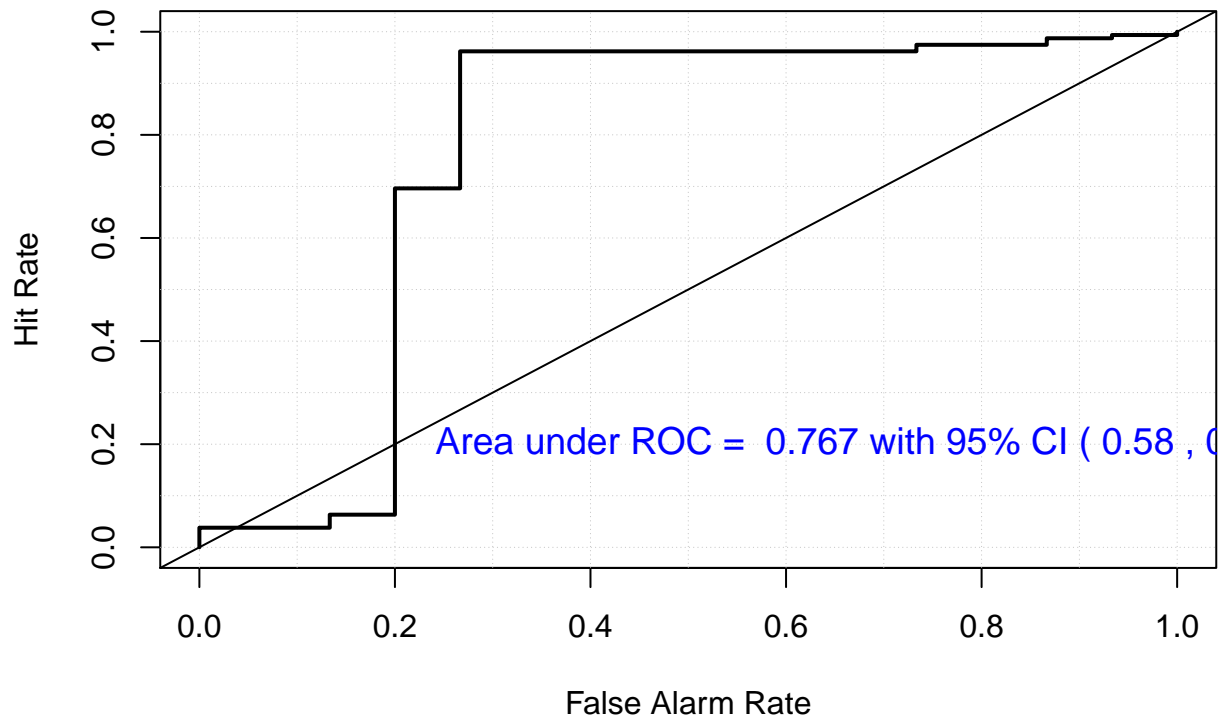
```
library(verification)
mod.nn2 <- verify(obs = yobs, pred = ypred_nn2_scale)
```

```
## If baseline is not included, baseline values will be calculated from the sample obs.
```

```
roc.plot(mod.nn2, plot.thres=NULL)
text(x=0.7, y=0.2, paste("Area under ROC = ", round(nn2_AUC$cvAUC, digits = 3), "with 95% CI (",
```

```
nn2_auc.ci[1], ",", nn2_auc.ci[2], ").", sep = " "), col="blue", cex =
```

ROC Curve



```
confusionMatrix(factor(ypred_nn22), factor(test$outcome))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    0    1
```

```
##           0  11   6
```

```
##           1   4 152
```

```
##
```

```
##           Accuracy : 0.9421965
```

```
##           95% CI : (0.8962682, 0.9719361)
```

```
##      No Information Rate : 0.9132948
```

```
##      P-Value [Acc > NIR] : 0.1076656
```

```
##
```

```
##           Kappa : 0.6557899
```

```
##  McNemar's Test P-Value : 0.7518296
```

```
##
```

```
##           Sensitivity : 0.7333333
```

```
##           Specificity : 0.9620253
```

```
##      Pos Pred Value : 0.6470588
```

```
##      Neg Pred Value : 0.9743589
```

```
##           Prevalence : 0.0867052
```

```
##      Detection Rate : 0.0635838
```

```
##      Detection Prevalence : 0.09826590
##      Balanced Accuracy   : 0.84767932
##
##      'Positive' Class    : 0
##
```

7 Support Vector Machine

I used five types of SVM linear techniques such as SVM linear, svm linear with grid, svm radial, svm radial grid, and Radial Basis kernel "Gaussian".

7.1 SVM-Liner

```
##### SVM Linear #####
## SUPPORT VECTOR MACHINE MODEL
library(caret)
set.seed(1492)
# Setup for cross validation
ctrl <- trainControl(method="repeatedcv", # 10fold cross validation
                     number =10)

set.seed(3233)

svm_Linear <- train(factor(outcome)~., data = train, method = "svmLinear",
                   trControl=ctrl,
                   tuneLength = 5) #5 arbitrary values of cost function C

svm_Linear

## Support Vector Machines with Linear Kernel
##
## 367 samples
## 18 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 331, 330, 330, 330, 330, 330, ...
## Resampling results:
##
##      Accuracy      Kappa
##      0.9425675676  0.553990734
##
## Tuning parameter 'C' was held constant at a value of 1
```

The confusion matrix produces the sensitivity, specificity, accuracy, and confidence interval etc.

```
svm_pred <- predict(svm_Linear, newdata = test);svm_pred

##      [1] 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1
##     [36] 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 0 1 1 1 1 1 1 1
##     [71] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
##    [106] 1 1 1 0 1 1 1 1 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1
##    [141] 1 1 0 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1
## Levels: 0 1
```

```
confusionMatrix(svm_pred, factor(test$outcome))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    0    1
```

```
##           0    9    5
```

```
##           1    6 153
```

```
##
```

```
##           Accuracy : 0.9364162
```

```
##           95% CI : (0.8890805, 0.9678347)
```

```
##      No Information Rate : 0.9132948
```

```
##      P-Value [Acc > NIR] : 0.1728098
```

```
##
```

```
##           Kappa : 0.5860344
```

```
##      Mcnemar's Test P-Value : 1.0000000
```

```
##
```

```
##           Sensitivity : 0.60000000
```

```
##           Specificity : 0.96835443
```

```
##           Pos Pred Value : 0.64285714
```

```
##           Neg Pred Value : 0.96226415
```

```
##           Prevalence : 0.08670520
```

```
##           Detection Rate : 0.05202312
```

```
##      Detection Prevalence : 0.08092486
```

```
##           Balanced Accuracy : 0.78417722
```

```
##
```

```
##           'Positive' Class : 0
```

```
##
```

```
(miss.rate_f <- mean(yobs != svm_pred))
```

```
## [1] 0.06358381503
```

```
svm_pred <- as.numeric(as.character(svm_pred))
```

```
MSE.f <- mean((yobs-svm_pred)^2)
```

```
MSE.f
```

```
## [1] 0.06358381503
```

Plotting ROC curve of the fit.best model.

```

library(cvAUC)
svm_lin_AUC <- ci.cvAUC(predictions = svm_pred, labels = yobs, folds=1:NROW(test), confidence = 0.95)
svm_lin_AUC

## $cvAUC
## [1] 0.7841772152
##
## $se
## [1] 0.1344847533
##
## $ci
## [1] 0.5205919423 1.0000000000
##
## $confidence
## [1] 0.95


(svm_lin_auc.ci <- round(svm_lin_AUC$ci, digits = 3))

## [1] 0.521 1.000

library(verification)
mod.svm_lin <- verify(obs = yobs, pred = svm_pred)

## If baseline is not included, baseline values will be calculated from the sample obs.
roc.plot(mod.svm_lin, plot.thres=NULL)
text(x=0.7, y=0.2, paste("Area under ROC = ", round(svm_lin_AUC$cvAUC, digits = 3), "with 95% CI [",
                        svm_lin_auc.ci[1], ",", svm_lin_auc.ci[2], "].", sep = " "), col="blue")

```



7.2 SVM-Linear-grid

```
##### C value in Linearclassifier #####
grid <- expand.grid(C = c(0.01, 0.05, 0.1, 0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2,5))
set.seed(3233)
svm_Linear_Grid <- train(factor(outcome) ~., data = train, method = "svmLinear",
                        trControl=ctrl,
                        tuneGrid = grid,
                        tuneLength = 5)

test_pred_grid <- predict(svm_Linear_Grid, newdata = test)
confusionMatrix(test_pred_grid, factor(test$outcome))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    0    1
```

```
##           0    7    2
```

```
##           1    8 156
```

```
##
```

```
##           Accuracy : 0.9421965
```

```
##           95% CI : (0.8962682, 0.9719361)
```

```
##           No Information Rate : 0.9132948
```



```
##      P-Value [Acc > NIR] : 0.1076656
##
##              Kappa : 0.5543534
##  McNemar's Test P-Value : 0.1138463
##
##      Sensitivity : 0.4666667
##      Specificity : 0.98734177
##      Pos Pred Value : 0.77777778
##      Neg Pred Value : 0.95121951
##      Prevalence : 0.08670520
##      Detection Rate : 0.04046243
##      Detection Prevalence : 0.05202312
##      Balanced Accuracy : 0.72700422
##
##      'Positive' Class : 0
##
```

```
(miss.rate_g <- mean(yobs != test_pred_grid))
```

```
## [1] 0.05780346821
```

```
test_pred_grid <- as.numeric(as.character(test_pred_grid))
MSE.g <- mean((yobs-test_pred_grid)^2)
MSE.g
```

```
## [1] 0.05780346821
```

Plotting ROC curve of the fit.best model.


```
library(cvAUC)
svm_lin_g_AUC <- ci.cvAUC(predictions = test_pred_grid, labels = yobs, folds=1:NROW(test), con=
(svm_lin_g_auc.ci <- round(svm_lin_g_AUC$ci, digits = 3))
```

```
## [1] 0.441 1.000
```

```
library(verification)
mod.svm_lin_g <- verify(obs = yobs, pred = test_pred_grid)
```


If baseline is not included, baseline values will be calculated from the sample obs.

```
roc.plot(mod.svm_lin_g, plot.thres=NULL)
text(x=0.7, y=0.2, paste("Area under ROC = ", round(svm_lin_g_AUC$cvAUC, digits = 3), "with 95%",
                        svm_lin_g_auc.ci[1], ",", svm_lin_g_auc.ci[2], ").", sep = " "), col=
```



7.3 SVM-Radial

```
#####SVM Classifier using Non-Linear Kernel #####  
set.seed(3233)  
svm_Radial <- train(factor(outcome) ~., data = train, method = "svmRadial", #radial kernel  
                    trControl=ctrl,  
                    tuneLength = 10)  
plot(svm_Radial)
```



```
test_pred_Radial <- predict(svm_Radial, newdata = test)
confusionMatrix(test_pred_Radial, factor(test$outcome))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    0    1
```

```
##           0    4    1
```

```
##           1   11  157
```

```
##
```

```
##           Accuracy : 0.9306358
```

```
##           95% CI : (0.88197, 0.9636472)
```

```
## No Information Rate : 0.9132948
```

```
## P-Value [Acc > NIR] : 0.256301653
```

```
##
```

```
##           Kappa : 0.3728097
```

```
## Mcnemar's Test P-Value : 0.009374768
```

```
##
```

```
##           Sensitivity : 0.2666667
```

```
##           Specificity : 0.99367089
```

```
## Pos Pred Value : 0.80000000
```

```
## Neg Pred Value : 0.93452381
```

```
## Prevalence : 0.08670520
```

```
## Detection Rate : 0.02312139
```

```
## Detection Prevalence : 0.02890173
```

```
## Balanced Accuracy : 0.63016878
```

```
##
##      'Positive' Class : 0
##
(miss.rate_h <- mean(yobs != test_pred_Radial))

## [1] 0.06936416185

test_pred_Radial <- as.numeric(as.character(test_pred_Radial))
MSE.h <- mean((yobs-test_pred_Radial)^2)
MSE.h

## [1] 0.06936416185

Plotting ROC curve of the best model.
library(cvAUC)
svm_rad_AUC <- ci.cvAUC(predictions = test_pred_Radial, labels = yobs, folds=1:NROW(test), con
svm_rad_AUC


## $cvAUC
## [1] 0.6301687764
##
## $se
## [1] 0.1503673744
##
## $ci
## [1] 0.3354541382 0.9248834146
##
## $confidence
## [1] 0.95

(svm_rad_auc.ci <- round(svm_rad_AUC$ci, digits = 3))

## [1] 0.335 0.925

library(verification)
mod.rad_h <- verify(obs = yobs, pred = test_pred_Radial)

## If baseline is not included, baseline values will be calculated from the sample obs.
roc.plot(mod.rad_h, plot.thres=NULL)
text(x=0.7, y=0.2, paste("Area under ROC = ", round(svm_rad_AUC$cvAUC, digits = 3), "with 95% (",
                        svm_rad_auc.ci[1], ",", svm_rad_auc.ci[2], ").", sep = " "), col="blue")
```



7.4 SVM-Radial-grid

```
grid_radial <- expand.grid(sigma = c(0.01, 0.02, 0.025, 0.03,
                                     0.05, 0.09, 0.1, 0.25, 0.5,
                                     0.75, 0.9), C = c(0.01, 0.05, 0.1, 0.25, 0.5, 0.75,
                                     1, 1.5, 2, 5 ))
```

```
set.seed(3233)
svm_Radial_Grid <- train(factor(outcome) ~ ., data = train, method = "svmRadial",
                        trControl=ctrl,
                        tuneGrid = grid_radial,
                        tuneLength = 9)
```

```
svm_Radial_Grid
```

```
## Support Vector Machines with Radial Basis Function Kernel
##
## 367 samples
## 18 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 1 times)
## Summary of sample sizes: 331, 330, 330, 330, 330, 330, ...
## Resampling results across tuning parameters:
```

```

##
##  sigma  C      Accuracy      Kappa
##  0.010  0.01  0.9155405405  0.000000000000
##  0.010  0.05  0.9155405405  0.000000000000
##  0.010  0.10  0.9155405405  0.000000000000
##  0.010  0.25  0.9155405405  0.000000000000
##  0.010  0.50  0.9155405405  0.000000000000
##  0.010  0.75  0.9155405405  0.000000000000
##  0.010  1.00  0.9155405405  0.000000000000
##  0.010  1.50  0.9155405405  0.000000000000
##  0.010  2.00  0.9155405405  0.000000000000
##  0.010  5.00  0.9590090090  0.63781707450
##  0.020  0.01  0.9155405405  0.000000000000
##  0.020  0.05  0.9155405405  0.000000000000
##  0.020  0.10  0.9155405405  0.000000000000
##  0.020  0.25  0.9155405405  0.000000000000
##  0.020  0.50  0.9155405405  0.000000000000
##  0.020  0.75  0.9155405405  0.000000000000
##  0.020  1.00  0.9155405405  0.000000000000
##  0.020  1.50  0.9209459459  0.09577464789
##  0.020  2.00  0.9373123123  0.33813425276
##  0.020  5.00  0.9563063063  0.63138650756
##  0.025  0.01  0.9155405405  0.000000000000
##  0.025  0.05  0.9155405405  0.000000000000
##  0.025  0.10  0.9155405405  0.000000000000
##  0.025  0.25  0.9155405405  0.000000000000
##  0.025  0.50  0.9155405405  0.000000000000
##  0.025  0.75  0.9155405405  0.000000000000
##  0.025  1.00  0.9155405405  0.000000000000
##  0.025  1.50  0.9292042042  0.21163421205
##  0.025  2.00  0.9454204204  0.44747236234
##  0.025  5.00  0.9480480480  0.58577562184
##  0.030  0.01  0.9155405405  0.000000000000
##  0.030  0.05  0.9155405405  0.000000000000
##  0.030  0.10  0.9155405405  0.000000000000
##  0.030  0.25  0.9155405405  0.000000000000
##  0.030  0.50  0.9155405405  0.000000000000
##  0.030  0.75  0.9155405405  0.000000000000
##  0.030  1.00  0.9155405405  0.000000000000
##  0.030  1.50  0.9346096096  0.29024692881
##  0.030  2.00  0.9399399399  0.38464997346
##  0.030  5.00  0.9424924925  0.53727195159
##  0.050  0.01  0.9155405405  0.000000000000
##  0.050  0.05  0.9155405405  0.000000000000
##  0.050  0.10  0.9155405405  0.000000000000
##  0.050  0.25  0.9155405405  0.000000000000
##  0.050  0.50  0.9155405405  0.000000000000
##  0.050  0.75  0.9155405405  0.000000000000

```


##	0.050	1.00	0.9182432432	0.04788732394
##	0.050	1.50	0.9400150150	0.36885964558
##	0.050	2.00	0.9454204204	0.45311146009
##	0.050	5.00	0.9397147147	0.49095405943
##	0.090	0.01	0.9155405405	0.00000000000
##	0.090	0.05	0.9155405405	0.00000000000
##	0.090	0.10	0.9155405405	0.00000000000
##	0.090	0.25	0.9155405405	0.00000000000
##	0.090	0.50	0.9155405405	0.00000000000
##	0.090	0.75	0.9155405405	0.00000000000
##	0.090	1.00	0.9182432432	0.04788732394
##	0.090	1.50	0.9265015015	0.17434607646
##	0.090	2.00	0.9318318318	0.30306103385
##	0.090	5.00	0.9373123123	0.35521499810
##	0.100	0.01	0.9155405405	0.00000000000
##	0.100	0.05	0.9155405405	0.00000000000
##	0.100	0.10	0.9155405405	0.00000000000
##	0.100	0.25	0.9155405405	0.00000000000
##	0.100	0.50	0.9155405405	0.00000000000
##	0.100	0.75	0.9155405405	0.00000000000
##	0.100	1.00	0.9155405405	0.00000000000
##	0.100	1.50	0.9265015015	0.17434607646
##	0.100	2.00	0.9292042042	0.22223340040
##	0.100	5.00	0.9264264264	0.21788557432
##	0.250	0.01	0.9155405405	0.00000000000
##	0.250	0.05	0.9155405405	0.00000000000
##	0.250	0.10	0.9155405405	0.00000000000
##	0.250	0.25	0.9155405405	0.00000000000
##	0.250	0.50	0.9155405405	0.00000000000
##	0.250	0.75	0.9155405405	0.00000000000
##	0.250	1.00	0.9155405405	0.00000000000
##	0.250	1.50	0.9155405405	0.00000000000
##	0.250	2.00	0.9155405405	0.00000000000
##	0.250	5.00	0.9155405405	0.00000000000
##	0.500	0.01	0.9155405405	0.00000000000
##	0.500	0.05	0.9155405405	0.00000000000
##	0.500	0.10	0.9155405405	0.00000000000
##	0.500	0.25	0.9155405405	0.00000000000
##	0.500	0.50	0.9155405405	0.00000000000
##	0.500	0.75	0.9155405405	0.00000000000
##	0.500	1.00	0.9155405405	0.00000000000
##	0.500	1.50	0.9155405405	0.00000000000
##	0.500	2.00	0.9155405405	0.00000000000
##	0.500	5.00	0.9155405405	0.00000000000
##	0.750	0.01	0.9155405405	0.00000000000
##	0.750	0.05	0.9155405405	0.00000000000
##	0.750	0.10	0.9155405405	0.00000000000
##	0.750	0.25	0.9155405405	0.00000000000

```
## 0.750 0.50 0.9155405405 0.00000000000
## 0.750 0.75 0.9155405405 0.00000000000
## 0.750 1.00 0.9155405405 0.00000000000
## 0.750 1.50 0.9155405405 0.00000000000
## 0.750 2.00 0.9155405405 0.00000000000
## 0.750 5.00 0.9155405405 0.00000000000
## 0.900 0.01 0.9155405405 0.00000000000
## 0.900 0.05 0.9155405405 0.00000000000
## 0.900 0.10 0.9155405405 0.00000000000
## 0.900 0.25 0.9155405405 0.00000000000
## 0.900 0.50 0.9155405405 0.00000000000
## 0.900 0.75 0.9155405405 0.00000000000
## 0.900 1.00 0.9155405405 0.00000000000
## 0.900 1.50 0.9155405405 0.00000000000
## 0.900 2.00 0.9155405405 0.00000000000
## 0.900 5.00 0.9155405405 0.00000000000
##
```

```
## Accuracy was used to select the optimal model using the largest value.
```

```
## The final values used for the model were sigma = 0.01 and C = 5.
```

```
plot(svm_Radial_Grid)
```



```
pred_Radial <- predict(svm_Radial_Grid, newdata = test)
confusionMatrix(pred_Radial, factor(test$outcome))
```

```
## Confusion Matrix and Statistics
```



```
##
##           Reference
## Prediction    0    1
##           0    5    1
##           1   10  157
##
##           Accuracy : 0.9364162
##           95% CI : (0.8890805, 0.9678347)
##    No Information Rate : 0.9132948
##    P-Value [Acc > NIR] : 0.17280985
##
##           Kappa : 0.448885
##    Mcnemar's Test P-Value : 0.01586133
##
##           Sensitivity : 0.3333333
##           Specificity : 0.99367089
##           Pos Pred Value : 0.8333333
##           Neg Pred Value : 0.94011976
##           Prevalence : 0.08670520
##           Detection Rate : 0.02890173
##    Detection Prevalence : 0.03468208
##           Balanced Accuracy : 0.66350211
##
##           'Positive' Class : 0
##
(miss.rate_i <- mean(yobs != pred_Radial))

## [1] 0.06358381503

pred_Radial <- as.numeric(as.character(pred_Radial))
MSE.i <- mean((yobs-pred_Radial)^2)
MSE.i

## [1] 0.06358381503

Plotting ROC curve of the fit.best model.

library(cvAUC)
svm3_AUC <- ci.cvAUC(predictions = pred_Radial, labels = yobs, folds=1:NROW(test), confidence = 0.95)
AUC

## $cvAUC
## [1] 0.9067510549
##
## $se
## [1] 0.0606345602
##
## $ci
## [1] 0.7879095006 1.0000000000
##
```

```
## $confidence
## [1] 0.95


(svm3_auc.ci <- round(svm3_AUC$ci, digits = 3))

## [1] 0.368 0.959

library(verification)
mod.svm3 <- verify(obs = yobs, pred = pred_Radial)

## If baseline is not included, baseline values will be calculated from the sample obs.
roc.plot(mod.svm3, plot.thres=NULL)
text(x=0.7, y=0.2, paste("Area under ROC = ", round(svm3_AUC$cvAUC, digits = 3), "with 95% CI",
                        svm3_auc.ci[1], ",", svm3_auc.ci[2], ").", sep = " "), col="blue", ce
```


ROC Curve



7.5 SVM: Radial Basis kernel "Gaussian"

```
library(kernlab)
library("e1071")
# USING ONLY PART OF THE DATA FOR TUNING FOR SHORTER COMPUTING TIME
tobj <- tune.svm(outcome ~ ., data = train, gamma = 10^(-6:-1), cost = 10^(-2:2), nrepeat=2,
                tunecontrol = tune.control(sampling = "cross", cross=10))

plot(tobj, transform.x = log10, xlab = expression(log[10](gamma)), ylab = "C")
```



```
bestGamma <- tobj$best.parameters[[1]]
bestC <- tobj$best.parameters[[2]]
```

```
svm_probab <- ksvm(outcome ~ ., data = train, type = "C-bsvc", kernel = "rbfdot",
  kpar = list(sigma = bestGamma), C = bestC,
  scaled=FALSE, prob.model = TRUE)
```

```
svm_probab_pred <- predict(svm_probab, type="probabilities", newdata=test)
svm_yhat <- svm_probab_pred[, 2] # EXTRACT PREDICTED PROB FOR spam
```

```
MSE.j <- mean((yobs-svm_yhat)^2)
MSE.j
```

```
## [1] 0.0450621658
```

```
svm_yhat_factor <- ifelse(svm_yhat>0.5, 1, 0)
(miss.rate_j <- mean(yobs != svm_yhat_factor))
```

```
## [1] 0.06936416185
```

```
library(cvAUC)
svmprob_AUC <- ci.cvAUC(predictions = svm_yhat, labels =yobs, folds=1:NROW(test), confidence =
(svmprob_auc.ci <- round(svmprob_AUC$ci, digits = 3))
```


```
## [1] 0.817 1.000
```

```
library(verification)
mod.svmprob <- verify(obs = yobs, pred = svm_yhat)
```

```
## If baseline is not included, baseline values will be calculated from the sample obs.
```

```
roc.plot(mod.svmprob, plot.thres=NULL)
text(x=0.7, y=0.2, paste("Area under ROC = ", round(svmprob_AUC$cvAUC, digits = 3), "with 95% CI (",
                        svmprob_auc.ci[1], ",", svmprob_auc.ci[2], ").", sep = " "), col="blue")
```

ROC Curve



```
confusionMatrix(factor(yobs), factor(svm_yhat_factor))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  0    1
```

```
##           0    6    9
```

```
##           1    3  155
```

```
##
```

```
##           Accuracy : 0.9306358
```

```
##           95% CI : (0.88197, 0.9636472)
```

```
##           No Information Rate : 0.9479769
```

```
##           P-Value [Acc > NIR] : 0.8816346
```

```
##
```

```
##           Kappa : 0.4652241
```

```
##           McNemar's Test P-Value : 0.1489147
```

```
##
```

```
##           Sensitivity : 0.6666667
```

```
##           Specificity : 0.94512195
```

```
##           Pos Pred Value : 0.40000000
```

```
##          Neg Pred Value : 0.98101266
##          Prevalence : 0.05202312
##          Detection Rate : 0.03468208
##          Detection Prevalence : 0.08670520
##          Balanced Accuracy : 0.80589431
##
##          'Positive' Class : 0
##
```

8 Comparison among the fitted models

We see the MSE of ten fitting models that I used in this project.

```
(tabulate <- data.frame(
  Methods = c("Logistic regression", "Random forest", "ANN (single layer)", "ANN(two layer)", "ANN(Three layers)", "SVM-linear", "SVM-linear-grid", "SVM-Radial", "SVM-Radial", "SVM-Basis-Kernel_Gaussian"),
  Mean.squared.error = c(MSE.a, MSE.b, MSE.c, MSE.d, MSE.e, MSE.f, MSE.g, MSE.h, MSE.i, MSE.j),
  Missclassification.rate = c(miss.rate_a, miss.rate_b, miss.rate_c, miss.rate_d, miss.rate_e, miss.rate_f, miss.rate_g, miss.rate_h, miss.rate_i, miss.rate_j))

##          Methods Mean.squared.error Missclassification.rate
## 1 Logistic regression      0.05003157544      0.07514450867
## 2 Random forest          0.06026632948      0.08670520231
## 3 ANN (single layer)      0.06982044344      0.06936416185
## 4 ANN(two layer)         0.08478119453      0.08670520231
## 5 ANN(Three layers)      0.06238087088      0.05780346821
## 6 SVM-linear             0.06358381503      0.06358381503
## 7 SVM-linear-grid        0.05780346821      0.05780346821
## 8 SVM-Radial             0.06936416185      0.06936416185
## 9 SVM-Radial             0.06358381503      0.06358381503
## 10 SVM-Basis-Kernel_Gaussian 0.04506216580      0.06936416185
```

Remarks: The SVM-Basis-Kernel_Gaussian gives the lowest predicted mean squared error. The ANN with three layers have the lowest missclassification rate.

Are under ROC curve:

```
(tabulate <- data.frame(
  Methods = c("Logistic regression", "Random forest", "ANN (single layer)", "ANN(two layer)", "ANN(Three layers)", "SVM-linear", "SVM-linear-grid", "SVM-Radial", "SVM-Radial", "SVM-Basis-Kernel_Gaussian"),
  Model.Accuracy.percentage = c(round(AUC$cvAUC, digits = 3), round(rf_AUC$cvAUC, digits = 3), round(net_AUC$cvAUC, digits = 3), round(nn_AUC$cvAUC, digits = 3), round(nn2_AUC$cvAUC, digits = 3), round(svm_lin_AUC$cvAUC, digits = 3), round(svm_lin_grid_AUC$cvAUC, digits = 3), round(svm_radial_AUC$cvAUC, digits = 3), round(svm_radial_AUC$cvAUC, digits = 3), round(svm_basis_kernel_gaussian_AUC$cvAUC, digits = 3))

##          Methods Model.Accuracy.percentage
## 1 Logistic regression      0.907
## 2 Random forest          0.856
## 3 ANN (single layer)      0.951
## 4 ANN(two layer)         0.912
## 5 ANN(Three layers)      0.767
## 6 SVM-linear             0.784
## 7 SVM-linear-grid        0.727
```

## 8	SVM-Radial	0.630
## 9	SVM-Radial	0.664
## 10	SVM-Basis-Kernel_Gaussian	0.919

Remarks: ANN single layer gives the most accuracy. At the same time, SVM-Basis-Kernel_Gaussian, logistic regression, ANN with two layers, Random forest perform very well as well.

ROC curve:

```
par(mfrow=c(3, 2))
roc.plot(mod.glm, plot.thres=NULL, main=" ROC-Logistic Regression")
text(x=0.7, y=0.2, paste("Area under ROC = ", round(AUC$cvAUC, digits = 3), "with 95% CI (",
                        auc.ci[1], ",", auc.ci[2], ").", sep = " "), col="blue", cex = 1.2)

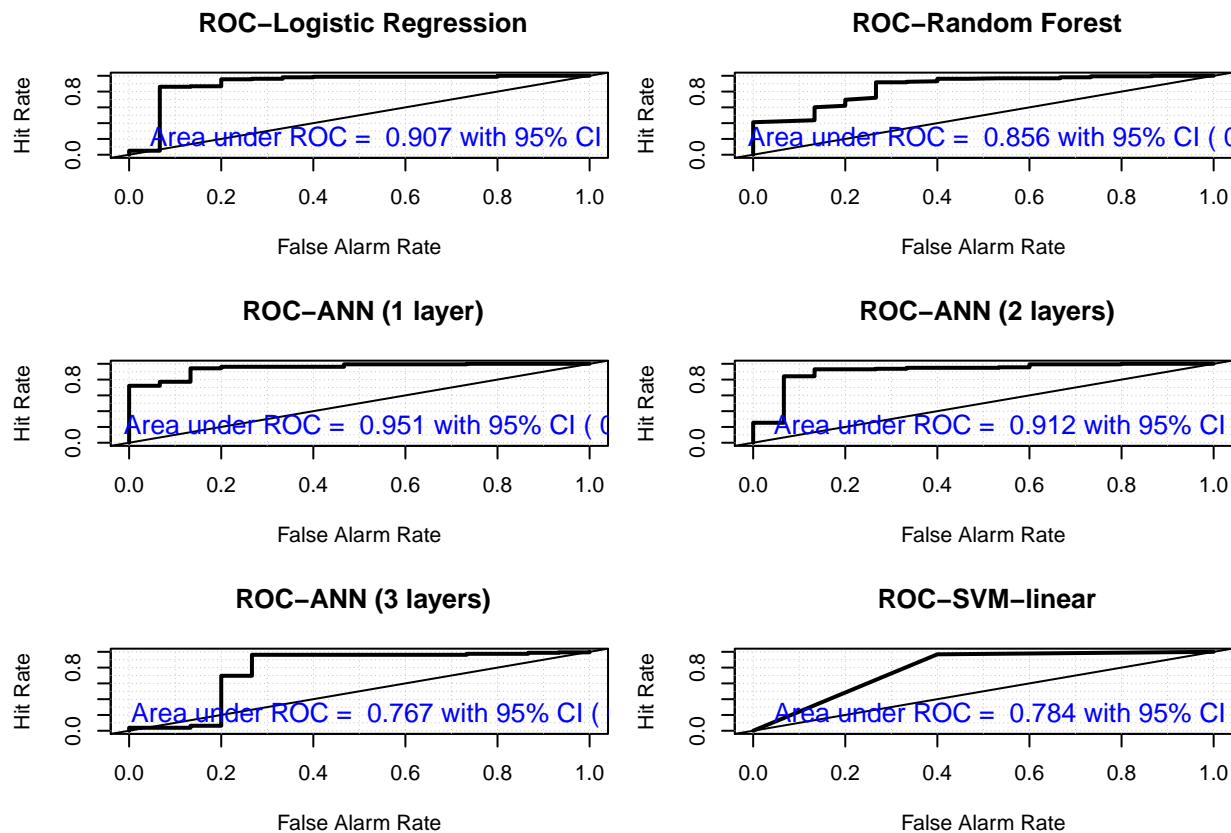
roc.plot(mod.rf, plot.thres=NULL, main=" ROC-Random Forest")
text(x=0.7, y=0.2, paste("Area under ROC = ", round(rf_AUC$cvAUC, digits = 3), "with 95% CI (",
                        rf_auc.ci[1], ",", rf_auc.ci[2], ").", sep = " "), col="blue", cex = 1.2)

roc.plot(mod.net1, plot.thres = NULL, main=" ROC-ANN (1 layer)")
text(x=0.7, y=0.2, paste("Area under ROC = ", round(net_AUC$cvAUC, digits = 3), "with 95% CI (",
                        net_auc.ci[1], ",", net_auc.ci[2], ").", sep = " "), col="blue", cex = 1.2)

roc.plot(mod.nn, plot.thres = NULL, main=" ROC-ANN (2 layers)")
text(x=0.7, y=0.2, paste("Area under ROC = ", round(nn_AUC$cvAUC, digits = 3), "with 95% CI (",
                        nn_auc.ci[1], ",", nn_auc.ci[2], ").", sep = " "), col="blue", cex = 1.2)

roc.plot(mod.nn2, plot.thres = NULL, main=" ROC-ANN (3 layers)")
text(x=0.7, y=0.2, paste("Area under ROC = ", round(nn2_AUC$cvAUC, digits = 3), "with 95% CI (",
                        nn2_auc.ci[1], ",", nn2_auc.ci[2], ").", sep = " "), col="blue", cex = 1.2)

roc.plot(mod.svm_lin, plot.thres=NULL, main=" ROC-SVM-linear")
text(x=0.7, y=0.2, paste("Area under ROC = ", round(svm_lin_AUC$cvAUC, digits = 3), "with 95% CI (",
                        svm_lin_auc.ci[1], ",", svm_lin_auc.ci[2], ").", sep = " "), col="blue", cex = 1.2)
```

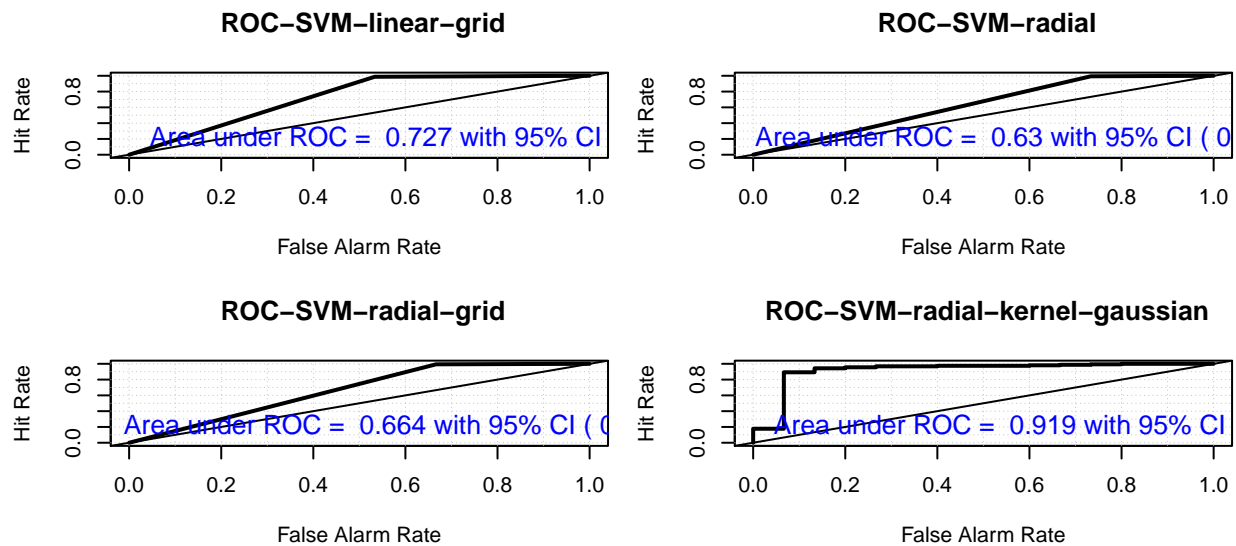


```
roc.plot(mod.svm_lin_g, plot.thres=NULL, main=" ROC-SVM-linear-grid")
text(x=0.7, y=0.2, paste("Area under ROC = ", round(svm_lin_g_AUC$cvAUC, digits = 3), "with 95% CI (",
                        svm_lin_g_auc.ci[1], ",", svm_lin_g_auc.ci[2], ").", sep = " "), col="blue")

roc.plot(mod.rad_h, plot.thres = NULL, main=" ROC-SVM-radial")
text(x=0.7, y=0.2, paste("Area under ROC = ", round(svm_rad_AUC$cvAUC, digits = 3), "with 95% CI (",
                        svm_rad_auc.ci[1], ",", svm_rad_auc.ci[2], ").", sep = " "), col="blue")

roc.plot(mod.svm3, plot.thres = NULL, main=" ROC-SVM-radial-grid")
text(x=0.7, y=0.2, paste("Area under ROC = ", round(svm3_AUC$cvAUC, digits = 3), "with 95% CI (",
                        svm3_auc.ci[1], ",", svm3_auc.ci[2], ").", sep = " "), col="blue", cex=1.2)

roc.plot(mod.svmprob, plot.thres = NULL, main=" ROC-SVM-radial-kernel-gaussian")
text(x=0.7, y=0.2, paste("Area under ROC = ", round(svmprob_AUC$cvAUC, digits = 3), "with 95% CI (",
                        svmprob_auc.ci[1], ",", svmprob_auc.ci[2], ").", sep = " "), col="blue", cex=1.2)
```



Remarks: From the ROC curve, we see that ANN single layer gives the most accuracy. At the same time, SVM-Basis-Kernel_Gaussian, Logistic Regression, ANN with two layers, Random Forest perform very well as well.